

# Task 0: Fashion MNIST classification in Pytorch (10 points)

The goal of this task is to get you familiar with [Pytorch](https://pytorch.org/) (<https://pytorch.org/>), teach you to debug your models, and give you a general understanding of deep learning and computer vision work-flows.

[Fashion MNIST](https://github.com/zalandoresearch/fashion-mnist) (<https://github.com/zalandoresearch/fashion-mnist>) is a dataset of [Zalando's](https://jobs.zalando.com/tech/) (<https://jobs.zalando.com/tech/>) article images — consisting of 70,000 grayscale images in 10 categories. Each example is a 28x28 grayscale image, associated with a label from 10 classes. ‘Fashion- MNIST’ is intended to serve as a direct **drop-in replacement** for the original [MNIST](http://yann.lecun.com/exdb/mnist/) (<http://yann.lecun.com/exdb/mnist/>) dataset — often used as the “Hello, World” of machine learning programs for computer vision. It shares the same image size and structure of training and testing splits. We will use 60,000 images to train the network and 10,000 images to evaluate how accurately the network learned to classify images.

```
In [1]: # installation directions can be found on pytorch's webpage
import torch
import torch.nn as nn
import torch.nn.functional as F
from torchvision import datasets, transforms
import matplotlib.pyplot as plt
%matplotlib inline

# import our network module from simple_cnn.py
from simple_cnn import SimpleCNN          # be sure to modify or you may have to restart kernel!
```

Usually you'll parse arguments using `argparse` (or similar library) but we can simply use a stand-in object for ipython notebooks. Furthermore, PyTorch can do computations on NVidia GPU s or on normal CPU s. You can configure the setting using the `device` variable.

```
In [2]: class ARGS(object):
    # input batch size for training
    batch_size = 64
    # input batch size for testing
    test_batch_size=1000
    # number of epochs to train for
    epochs = 14
    # Learning rate
    lr = 1.0
    # Learning rate step gamma
    gamma = 0.7
    # how many batches to wait before logging training status
    log_every = 100
    # how many batches to wait before evaluating model
    val_every = 100
    # set true if using GPU during training
    use_cuda = True

args = ARGS()
device = torch.device("cuda" if args.use_cuda else "cpu")
```

We define some basic testing and training code. The testing code prints out the average test loss and the training code ( `main` ) plots train/test losses and returns the final model.

```

In [3]: def test(model, device, test_loader):
    """Evaluate model on test dataset."""
    model.eval()
    test_loss = 0
    correct = 0
    with torch.no_grad():
        for data, target in test_loader:
            data, target = data.to(device), target.to(device)
            output = model(data)
            test_loss += F.cross_entropy(output, target, reduction='sum').item()
    # sum up batch loss
    pred = output.argmax(dim=1, keepdim=True) # get the index of the max log-probability
    correct += pred.eq(target.view_as(pred)).sum().item()

    test_loss /= len(test_loader.dataset)
    # print("TEST ACCURACY: ", 100. * correct / len(test_loader.dataset))
    print('\nTest set: Average loss: {:.4f}, Accuracy: {}/{} ({:.0f}%)\n'.format(
        test_loss, correct, len(test_loader.dataset),
        100. * correct / len(test_loader.dataset)))

    return test_loss, correct / len(test_loader.dataset)

def main():
    # 1. Load dataset and build dataloader
    train_loader = torch.utils.data.DataLoader(
        datasets.FashionMNIST('../data', train=True, download=True,
                              transform=transforms.Compose([
                                  transforms.ToTensor(),
                                  transforms.Normalize((0.1307,), (0.3081,))]))
    batch_size=args.batch_size, shuffle=True)
    test_loader = torch.utils.data.DataLoader(
        datasets.FashionMNIST('../data', train=False, transform=transforms.Compose([
            transforms.ToTensor(),
            transforms.Normalize((0.1307,), (0.3081,))]))
    batch_size=args.test_batch_size, shuffle=True)

    # 2. define the model, and optimizer.
    model = SimpleCNN().to(device)
    model.train()
    optimizer = torch.optim.Adam(model.parameters(), lr=args.lr)

    scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=1, gamma=
args.gamma)
    cnt = 0
    train_log = {'iter': [], 'loss': [], 'accuracy': []}
    test_log = {'iter': [], 'loss': [], 'accuracy': []}
    for epoch in range(args.epochs):
        for batch_idx, (data, target) in enumerate(train_loader):
            # Get a batch of data
            data, target = data.to(device), target.to(device)

```

```

optimizer.zero_grad()
# print("Size of data epoch data input: ",data.size())
# Forward pass
output = model(data)
# Calculate the loss
loss = F.cross_entropy(output, target)
# Calculate gradient w.r.t the loss
loss.backward()
# Optimizer takes one step
optimizer.step()
# Log info
if cnt % args.log_every == 0:
    print('Train Epoch: {} [{}/{} ({:.0f}%)]\tLoss: {:.6f}'.format(
        epoch, cnt, len(train_loader.dataset),
        100. * batch_idx / len(train_loader), loss.item()))
train_log['iter'].append(cnt)
train_log['loss'].append(loss)
# TODO: calculate your train accuracy!
mod_pred = torch.argmax(output, dim=1)
correct = mod_pred==target
correct = correct.long()
correct = sum(correct)
train_acc = correct.float()/args.batch_size
train_log['accuracy'].append(train_acc)
# print("TRAINING ACCURACY: ", train_acc)
# Validation iteration
if cnt % args.val_every == 0:
    test_loss, test_acc = test(model, device, test_loader)
    test_log['iter'].append(cnt)
    test_log['loss'].append(test_loss)
    test_log['accuracy'].append(test_acc)
    model.train()
    cnt += 1
scheduler.step()
fig = plt.figure()
plt.plot(train_log['iter'], train_log['loss'], 'r', label='Training')
plt.plot(test_log['iter'], test_log['loss'], 'b', label='Testing')
plt.title('Loss')
plt.legend()
fig = plt.figure()
plt.plot(train_log['iter'], train_log['accuracy'], 'r', label='Training')
plt.plot(test_log['iter'], test_log['accuracy'], 'b', label='Testing')
plt.title('Accuracy')
plt.legend()
plt.show()
return model

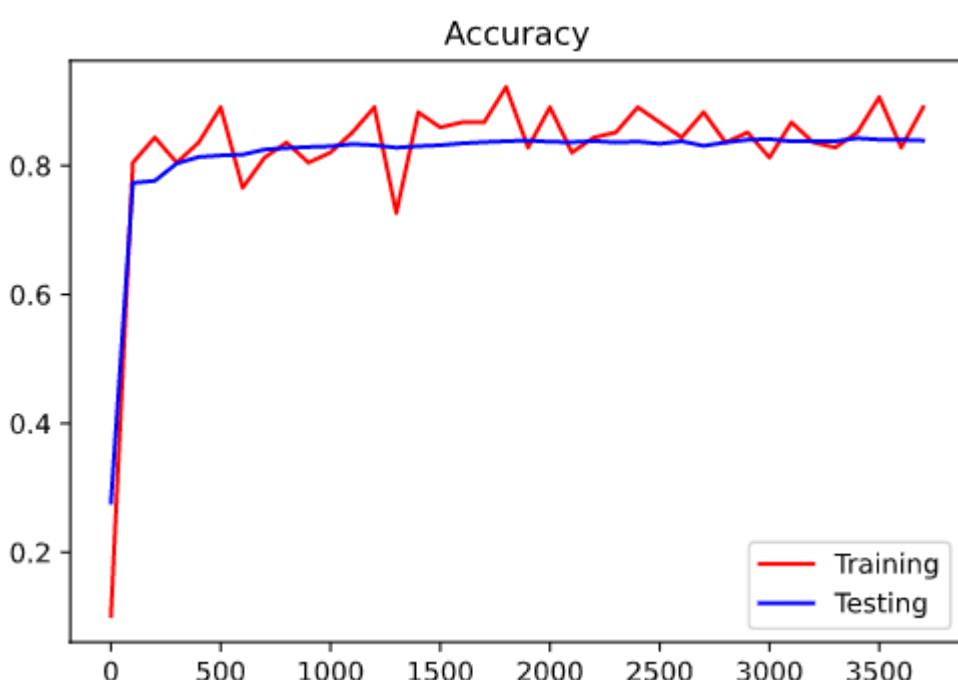
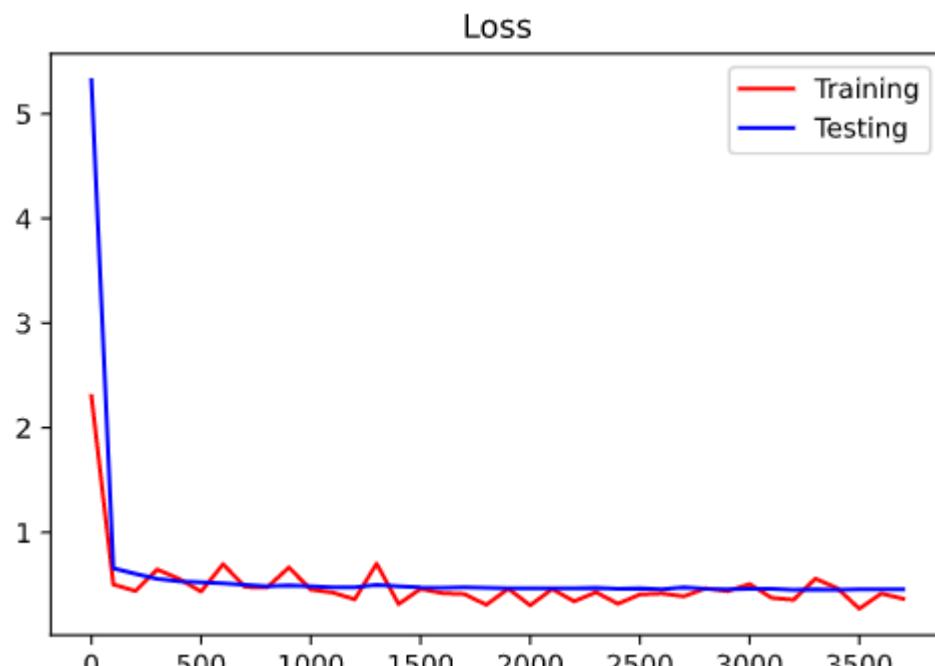
```

## 0.1 Bug Fix and Hyper-parameter search. (2pts)

Simply running `main` will result in a `RuntimeError` ! Check out `simple_cnn.py` and see if you can fix the bug. You may have to restart your ipython kernel for changes to reflect in the notebook. After that's done, be sure to fill in the TODOs in `main`.

Once you fix the bugs, you should be able to get a reasonable accuracy within 100 iterations just by tuning some hyper-parameter. Include the train/test plots of your best hyperparameter setting and comment on why you think these settings worked best. (you can complete this task on CPU)

YOUR ANSWER HERE



During the training period with the old hyper-parameters, one can notice that the accuracy of the test and training data was plateauing after a set number of epochs. As such, the number of epochs were reduced to 8. The batch size was increased to 128 to speed up the training period as well. The learning rate was significantly large so it was reduced down to 0.01. In the end, the accuracy of the training and testing sets were able to break the 80% mark.

```
In [4]: ##### FEEL FREE TO MODIFY args VARIABLE HERE OR ABOVE #####
# args.gamma = float('inf')
args.lr = 0.01 # lowered
args.epochs = 8
args.batch_size = 128

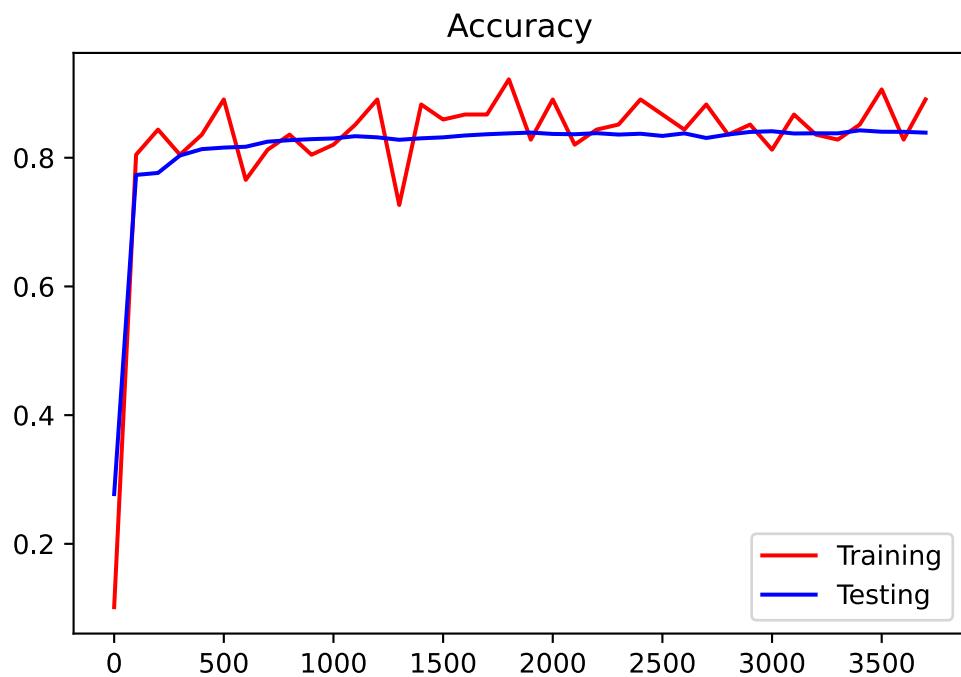
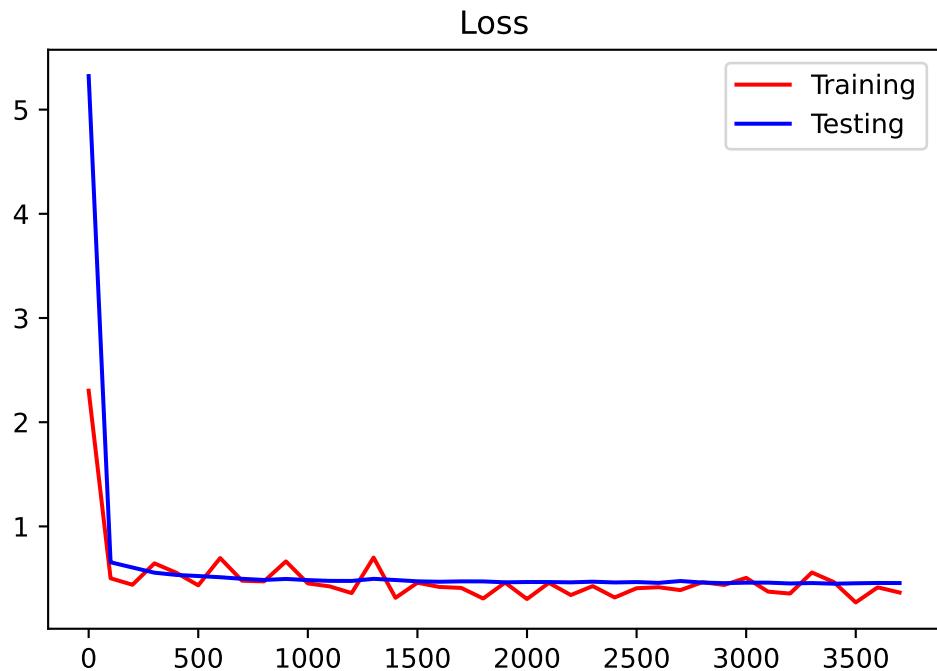
# DON'T CHANGE
# prints out arguments and runs main
for attr in dir(args):
    if '__' not in attr and attr != 'use_cuda':
        print('args.{0} = {1}'.format(attr, getattr(args, attr)))
print('\n\n')
model = main()
```

```
args.batch_size = 128
args.epochs = 8
args.gamma = 0.7
args.log_every = 100
args.lr = 0.01
args.test_batch_size = 1000
args.val_every = 100
```

```
Train Epoch: 0 [0/60000 (0%)] Loss: 2.302593
Test set: Average loss: 5.3219, Accuracy: 2773/10000 (28%)
Train Epoch: 0 [100/60000 (21%)] Loss: 0.501979
Test set: Average loss: 0.6558, Accuracy: 7733/10000 (77%)
Train Epoch: 0 [200/60000 (43%)] Loss: 0.439888
Test set: Average loss: 0.6055, Accuracy: 7764/10000 (78%)
Train Epoch: 0 [300/60000 (64%)] Loss: 0.645565
Test set: Average loss: 0.5559, Accuracy: 8035/10000 (80%)
Train Epoch: 0 [400/60000 (85%)] Loss: 0.555263
Test set: Average loss: 0.5344, Accuracy: 8134/10000 (81%)
Train Epoch: 1 [500/60000 (7%)] Loss: 0.434446
Test set: Average loss: 0.5238, Accuracy: 8158/10000 (82%)
Train Epoch: 1 [600/60000 (28%)] Loss: 0.696253
Test set: Average loss: 0.5123, Accuracy: 8171/10000 (82%)
Train Epoch: 1 [700/60000 (49%)] Loss: 0.478521
Test set: Average loss: 0.4971, Accuracy: 8250/10000 (82%)
Train Epoch: 1 [800/60000 (71%)] Loss: 0.472464
Test set: Average loss: 0.4861, Accuracy: 8273/10000 (83%)
Train Epoch: 1 [900/60000 (92%)] Loss: 0.664430
Test set: Average loss: 0.4959, Accuracy: 8289/10000 (83%)
Train Epoch: 2 [1000/60000 (13%)] Loss: 0.452996
Test set: Average loss: 0.4847, Accuracy: 8300/10000 (83%)
Train Epoch: 2 [1100/60000 (35%)] Loss: 0.424224
Test set: Average loss: 0.4772, Accuracy: 8335/10000 (83%)
```

Train Epoch: 2 [1200/60000 (56%)] Loss: 0.360117  
Test set: Average loss: 0.4767, Accuracy: 8317/10000 (83%)  
Train Epoch: 2 [1300/60000 (77%)] Loss: 0.701172  
Test set: Average loss: 0.4968, Accuracy: 8279/10000 (83%)  
Train Epoch: 2 [1400/60000 (99%)] Loss: 0.314604  
Test set: Average loss: 0.4855, Accuracy: 8302/10000 (83%)  
Train Epoch: 3 [1500/60000 (20%)] Loss: 0.459535  
Test set: Average loss: 0.4729, Accuracy: 8317/10000 (83%)  
Train Epoch: 3 [1600/60000 (41%)] Loss: 0.419015  
Test set: Average loss: 0.4688, Accuracy: 8346/10000 (83%)  
Train Epoch: 3 [1700/60000 (62%)] Loss: 0.409575  
Test set: Average loss: 0.4726, Accuracy: 8365/10000 (84%)  
Train Epoch: 3 [1800/60000 (84%)] Loss: 0.308018  
Test set: Average loss: 0.4721, Accuracy: 8379/10000 (84%)  
Train Epoch: 4 [1900/60000 (5%)] Loss: 0.460715  
Test set: Average loss: 0.4635, Accuracy: 8391/10000 (84%)  
Train Epoch: 4 [2000/60000 (26%)] Loss: 0.302965  
Test set: Average loss: 0.4663, Accuracy: 8370/10000 (84%)  
Train Epoch: 4 [2100/60000 (48%)] Loss: 0.458157  
Test set: Average loss: 0.4661, Accuracy: 8365/10000 (84%)  
Train Epoch: 4 [2200/60000 (69%)] Loss: 0.341236  
Test set: Average loss: 0.4631, Accuracy: 8382/10000 (84%)  
Train Epoch: 4 [2300/60000 (90%)] Loss: 0.427276  
Test set: Average loss: 0.4694, Accuracy: 8360/10000 (84%)  
Train Epoch: 5 [2400/60000 (12%)] Loss: 0.317789  
Test set: Average loss: 0.4621, Accuracy: 8373/10000 (84%)  
Train Epoch: 5 [2500/60000 (33%)] Loss: 0.406668  
Test set: Average loss: 0.4658, Accuracy: 8339/10000 (83%)

Train Epoch: 5 [2600/60000 (54%)] Loss: 0.416200  
Test set: Average loss: 0.4581, Accuracy: 8379/10000 (84%)  
Train Epoch: 5 [2700/60000 (76%)] Loss: 0.388463  
Test set: Average loss: 0.4759, Accuracy: 8307/10000 (83%)  
Train Epoch: 5 [2800/60000 (97%)] Loss: 0.464706  
Test set: Average loss: 0.4614, Accuracy: 8361/10000 (84%)  
Train Epoch: 6 [2900/60000 (18%)] Loss: 0.438095  
Test set: Average loss: 0.4554, Accuracy: 8401/10000 (84%)  
Train Epoch: 6 [3000/60000 (40%)] Loss: 0.506237  
Test set: Average loss: 0.4599, Accuracy: 8412/10000 (84%)  
Train Epoch: 6 [3100/60000 (61%)] Loss: 0.374118  
Test set: Average loss: 0.4604, Accuracy: 8377/10000 (84%)  
Train Epoch: 6 [3200/60000 (82%)] Loss: 0.355285  
Test set: Average loss: 0.4524, Accuracy: 8380/10000 (84%)  
Train Epoch: 7 [3300/60000 (4%)] Loss: 0.557764  
Test set: Average loss: 0.4565, Accuracy: 8379/10000 (84%)  
Train Epoch: 7 [3400/60000 (25%)] Loss: 0.465712  
Test set: Average loss: 0.4502, Accuracy: 8426/10000 (84%)  
Train Epoch: 7 [3500/60000 (46%)] Loss: 0.270384  
Test set: Average loss: 0.4542, Accuracy: 8404/10000 (84%)  
Train Epoch: 7 [3600/60000 (68%)] Loss: 0.415405  
Test set: Average loss: 0.4572, Accuracy: 8403/10000 (84%)  
Train Epoch: 7 [3700/60000 (89%)] Loss: 0.365243  
Test set: Average loss: 0.4563, Accuracy: 8390/10000 (84%)



## Play with parameters.(3pt)

How many trainable parameters does the trained model have?

Model has 454922 params

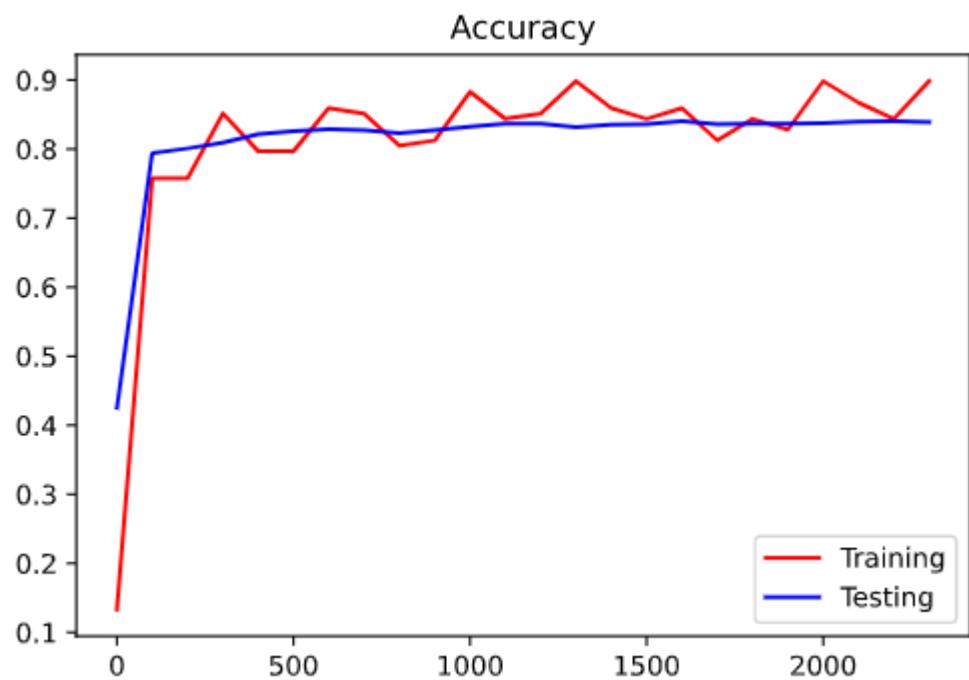
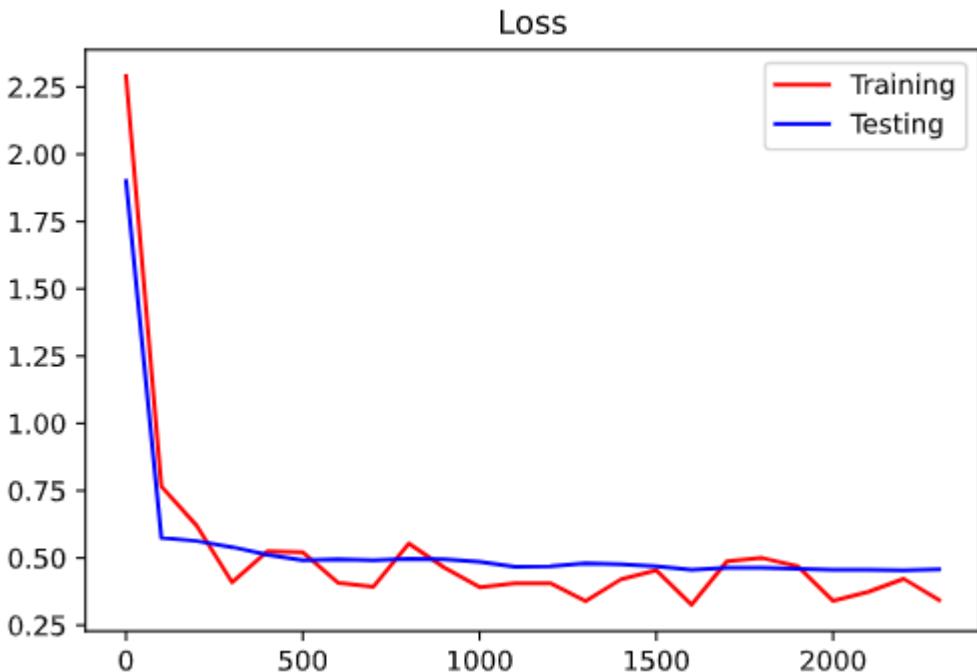
```
In [7]: def param_count(model):
    try:
        numParam = sum(iterVal.numel() for iterVal in model.parameters() if it
erVal.requires_grad)
        return numParam
    except:
        raise NotImplementedError
print('Model has {} params'.format(param_count(model)))
```

Model has 454922 params

## Deep Linear Networks?!? (5pt)

Until this point, there are no non-linearities in the SimpleCNN! (Your TAs were just as surprised as you are at the results.) Your next task is to modify the code to add non-linear activation layers, and train your model in full scale. Make sure to add non-linearities at **every** applicable layer.

Compute the loss and accuracy curves on train and test sets after 5 epochs.



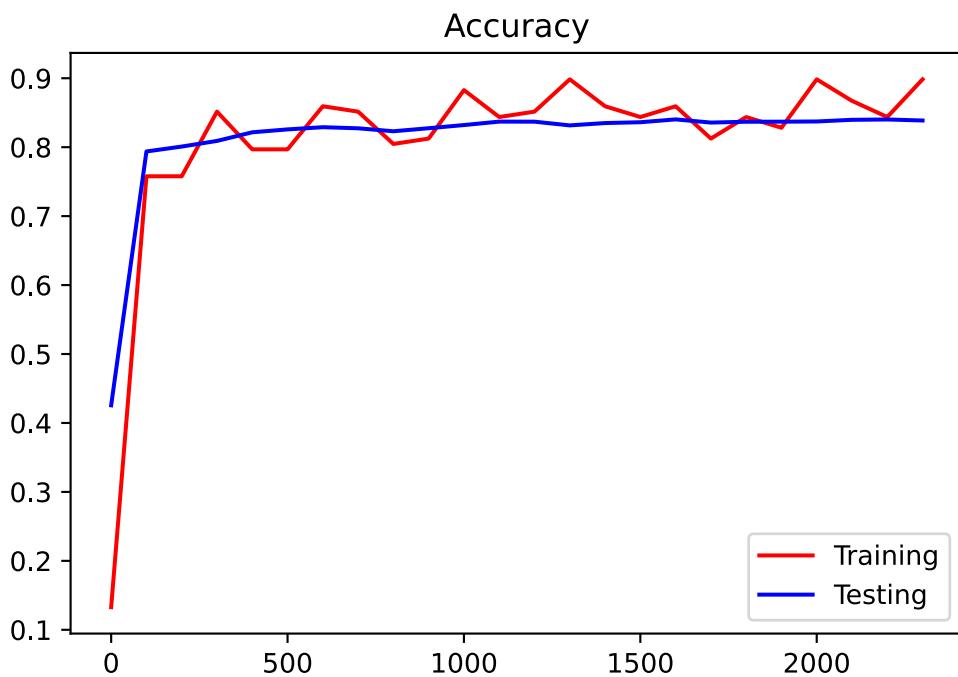
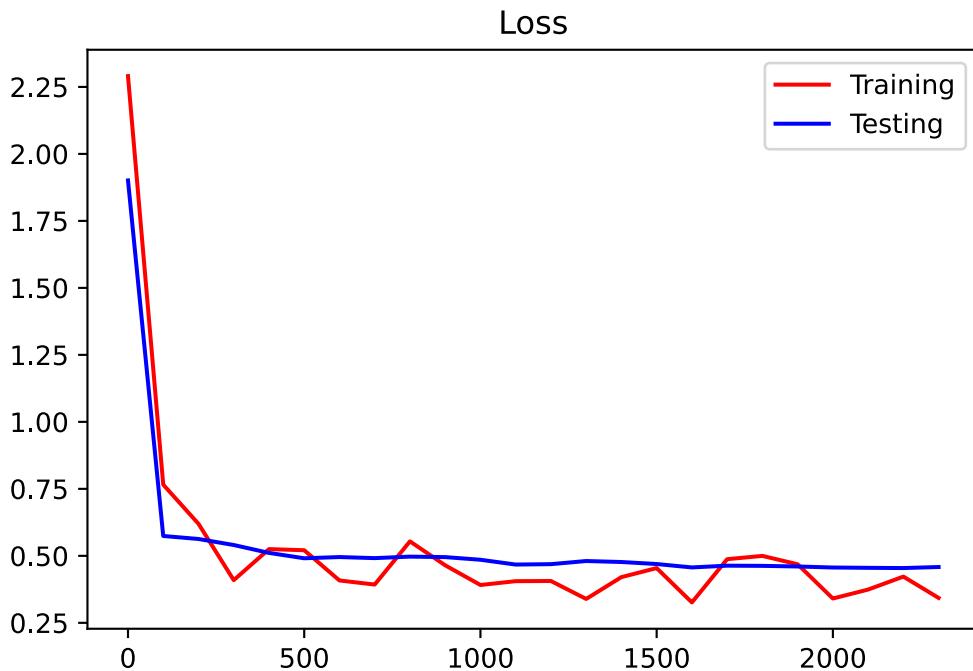


In [8]:

```
args.epochs = 5
args.lr = 1e-3
main()
```

Train Epoch: 0 [0/60000 (0%)] Loss: 2.290475  
Test set: Average loss: 1.9010, Accuracy: 4256/10000 (43%)  
Train Epoch: 0 [100/60000 (21%)] Loss: 0.765751  
Test set: Average loss: 0.5742, Accuracy: 7938/10000 (79%)  
Train Epoch: 0 [200/60000 (43%)] Loss: 0.619610  
Test set: Average loss: 0.5629, Accuracy: 8009/10000 (80%)  
Train Epoch: 0 [300/60000 (64%)] Loss: 0.409301  
Test set: Average loss: 0.5404, Accuracy: 8091/10000 (81%)  
Train Epoch: 0 [400/60000 (85%)] Loss: 0.525194  
Test set: Average loss: 0.5110, Accuracy: 8216/10000 (82%)  
Train Epoch: 1 [500/60000 (7%)] Loss: 0.520956  
Test set: Average loss: 0.4909, Accuracy: 8258/10000 (83%)  
Train Epoch: 1 [600/60000 (28%)] Loss: 0.408131  
Test set: Average loss: 0.4955, Accuracy: 8290/10000 (83%)  
Train Epoch: 1 [700/60000 (49%)] Loss: 0.392827  
Test set: Average loss: 0.4915, Accuracy: 8274/10000 (83%)  
Train Epoch: 1 [800/60000 (71%)] Loss: 0.554026  
Test set: Average loss: 0.4971, Accuracy: 8231/10000 (82%)  
Train Epoch: 1 [900/60000 (92%)] Loss: 0.464746  
Test set: Average loss: 0.4953, Accuracy: 8275/10000 (83%)  
Train Epoch: 2 [1000/60000 (13%)] Loss: 0.391170  
Test set: Average loss: 0.4854, Accuracy: 8321/10000 (83%)  
Train Epoch: 2 [1100/60000 (35%)] Loss: 0.405492  
Test set: Average loss: 0.4675, Accuracy: 8371/10000 (84%)  
Train Epoch: 2 [1200/60000 (56%)] Loss: 0.406192  
Test set: Average loss: 0.4690, Accuracy: 8370/10000 (84%)  
Train Epoch: 2 [1300/60000 (77%)] Loss: 0.339169  
Test set: Average loss: 0.4807, Accuracy: 8316/10000 (83%)  
Train Epoch: 2 [1400/60000 (99%)] Loss: 0.420615

Test set: Average loss: 0.4770, Accuracy: 8350/10000 (84%)  
Train Epoch: 3 [1500/60000 (20%)] Loss: 0.454342  
Test set: Average loss: 0.4695, Accuracy: 8362/10000 (84%)  
Train Epoch: 3 [1600/60000 (41%)] Loss: 0.326145  
Test set: Average loss: 0.4568, Accuracy: 8403/10000 (84%)  
Train Epoch: 3 [1700/60000 (62%)] Loss: 0.487575  
Test set: Average loss: 0.4633, Accuracy: 8357/10000 (84%)  
Train Epoch: 3 [1800/60000 (84%)] Loss: 0.499794  
Test set: Average loss: 0.4626, Accuracy: 8369/10000 (84%)  
Train Epoch: 4 [1900/60000 (5%)] Loss: 0.468573  
Test set: Average loss: 0.4603, Accuracy: 8370/10000 (84%)  
Train Epoch: 4 [2000/60000 (26%)] Loss: 0.340706  
Test set: Average loss: 0.4565, Accuracy: 8373/10000 (84%)  
Train Epoch: 4 [2100/60000 (48%)] Loss: 0.374197  
Test set: Average loss: 0.4553, Accuracy: 8397/10000 (84%)  
Train Epoch: 4 [2200/60000 (69%)] Loss: 0.422526  
Test set: Average loss: 0.4544, Accuracy: 8402/10000 (84%)  
Train Epoch: 4 [2300/60000 (90%)] Loss: 0.342963  
Test set: Average loss: 0.4583, Accuracy: 8387/10000 (84%)



**Out[8]:** SimpleCNN(  
    (conv1): Conv2d(1, 32, kernel\_size=(5, 5), stride=(1, 1), padding=(2, 2))  
    (conv2): Conv2d(32, 64, kernel\_size=(5, 5), stride=(1, 1), padding=(2, 2))  
    (pool1): AvgPool2d(kernel\_size=2, stride=2, padding=0)  
    (pool2): AvgPool2d(kernel\_size=2, stride=2, padding=0)  
    (fc1): Sequential(  
        (0): Linear(in\_features=3136, out\_features=128, bias=True)  
    )  
    (fc2): Sequential(  
        (0): Linear(in\_features=128, out\_features=10, bias=True)  
    )  
)

Where did you add your non-linearities?

**YOUR ANSWER HERE**

I added ReLU activation layers after each of the two convolution layers. I also added a ReLU activation after the first fully-connected layer. I did not apply a softmax at the second fully-connected layer because the cross\_entropy loss function defined in main() contains the softmax function.

Provide some insights on why the results was fairly good even without activation layers. (2 pts)

**YOUR ANSWER HERE**

The contents of the objects in the images of the dataset were presented in a standardized way. There was no occlusion or varying configurations of most of the objects in each class. For this reason, a model can generalize the dataset with good performance through the use of linear filters and logistic regression lines. The classification for shoes, for example, can be adequately characterized through a set of logistic regression lines along the edges of the shoe. The same thing may be said for jackets and pants. Differentiations between classes may be significant enough to be separated via logistic regression rather than a nonlinear function.

Another reason may be that the images of the same classification reside within the same sections of the frames compared to other class images. For example, trouser images usually occupy a rectangular area around the center of the image frame. Ankle boots usually occupy the lower-right diagonal half of the image. Given a linear network model with no activation functions, straight regression lines can be used to encompass common firing zones for certain classes. From this thought, nonlinear activations are not needed for the success of the model.



# Q1: Simple CNN network for PASCAL multi-label classification (20 points)

Now let's try to recognize some natural images. We provided some starter code for this task. The following steps will guide you through the process.

## 1.1 Setup the dataset

We start by modifying the code to read images from the PASCAL 2007 dataset. The important thing to note is that PASCAL can have multiple objects present in the same image. Hence, this is a multi-label classification problem, and will have to be tackled slightly differently.

First, download the data. `cd` to a location where you can store 0.5GB of images. Then run:

```
wget http://host.robots.ox.ac.uk/pascal/VOC/voc2007/VOCtrainval_06-Nov-2007.tar  
tar -xf VOCtrainval_06-Nov-2007.tar  
  
wget http://host.robots.ox.ac.uk/pascal/VOC/voc2007/VOCtest_06-Nov-2007.tar  
tar -xf VOCtest_06-Nov-2007.tar  
cd VOCdevkit/VOC2007/
```

## 1.2 Write a dataloader with data augmentation (5 pts)

**Dataloader** The first step is to write a [pytorch data loader](#)

([https://pytorch.org/tutorials/beginner/data\\_loading\\_tutorial.html](https://pytorch.org/tutorials/beginner/data_loading_tutorial.html)) which loads this PASCAL data. Complete the functions `preload_anno` and `__getitem__` in `voc_dataset.py`.

- **Hint:** Refer to the README in VOCdevkit to understand the structure and labeling.
- **Hint :** As the function docstring says, `__getitem__` takes as input the index, and returns a tuple - (`image`, `label`, `weight`) . The labels should be 1s for each object that is present in the image, and weights should be 1 for each label in the image, except those labeled as ambiguous (use the `difficult` attribute). All other values should be 0. For simplicity, resize all images to a canonical size.)

**Data Augmentation** Modify `__getitem__` to randomly *augment* each datapoint. Please describe what data augmentation you implement.

- **Hint:** Since we are training a model from scratch on this small dataset, it is important to perform basic data augmentation to avoid overfitting. Add random crops and left-right flips when training, and do a center crop when testing, etc. As for natural images, another common practice is to subtract the mean values of RGB images from ImageNet dataset. The mean values for RGB images are: [123.68, 116.78, 103.94] – sometimes, rescaling to [-1, 1] suffices.

**Note:** You should use data in ‘trainval’ for training and ‘test’ for testing, since PASCAL is a small dataset.

## DESCRIBE YOUR AUGMENTATION PIPELINE HERE\*\*

### Train Augmentations:

The following augmentations were performed on the training set:

- Rescaling to 28 x 28
- Random horizontal flipping
- Normalization based on a mean of (0.485, 0.456, 0.406) and a standard deviation of (0.229, 0.224, 0.225)

### Test Augmentations:

The following augmentations were performed on the test set:

- Rescaling to 28 x 28
- Random horizontal flipping
- Normalization based on a mean of (0.485, 0.456, 0.406) and a standard deviation of (0.229, 0.224, 0.225)

## 1.3 Measure Performance (5 pts)

To evaluate the trained model, we will use a standard metric for multi-label evaluation - [mean average precision \(mAP\)](http://scikit-learn.org/stable/modules/generated/sklearn.metrics.average_precision_score.html). Please implement `eval_dataset_map` in `utils.py` - this function will evaluate a model's map score using a given dataset object. You will need to make predictions on the given dataset with the model and call `compute_ap` to get average precision.

Please describe how to compute AP for each class(not mAP). **YOUR ANSWER HERE**

The average precision (AP) is computed by getting the ground truth, non-difficult classes and the correct prediction classes. The prediction class get a small decimal subtracted from it (perhaps for normalization). Then an average precision score is calculated between the two vectors and given a value for each class and stacked on top of each other as a vector. At the end of the function, the AP vector is returned out.

## 1.4 Let's Start Training! (5 pts)

Write the code for training and testing for multi-label classification in `trainer.py`. To start, you'll use the same model you used for Fashion MNIST (bad idea, but let's give it a shot).

Initialize a fresh model and optimizer. Then run your training code for 5 epochs and print the mAP on test set.

```
In [1]: import torch
import trainer
from utils import ARGS
from simple_cnn import SimpleCNN
from voc_dataset import VOCDataset

# create hyperparameter argument class
args = ARGS(epochs=5)
print(args)

args.batch_size = 32
args.device = cuda
args.epochs = 5
args.gamma = 0.7
args.log_every = 100
args.lr = 1.0
args.save_at_end = False
args.save_freq = 10
args.test_batch_size = 1000
args.val_every = 100
```

```
In [ ]: # initializes (your) naive model
model = SimpleCNN(num_classes=len(VOCDataset.CLASS_NAMES), inp_size=64, c_dim=3)
# initializes Adam optimizer and simple StepLR scheduler
optimizer = torch.optim.Adam(model.parameters(), lr=args.lr)
scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=1, gamma=args.gamma)
# trains model using your training code and reports test map
test_ap, test_map = trainer.train(args, model, optimizer, scheduler)
print('test map:', test_map)
```

[TensorBoard](https://www.tensorflow.org/guide/summaries_and_tensorboard) ([https://www.tensorflow.org/guide/summaries\\_and\\_tensorboard](https://www.tensorflow.org/guide/summaries_and_tensorboard)) is an awesome visualization tool. It was firstly integrated in [TensorFlow](https://www.tensorflow.org/) (<https://www.tensorflow.org/>) (~possibly the only useful tool TensorFlow provides~). It can be used to visualize training losses, network weights and other parameters.

To use TensorBoard in Pytorch, there are two options: [TensorBoard in Pytorch](https://pytorch.org/docs/stable/tensorboard.html) (<https://pytorch.org/docs/stable/tensorboard.html>) (for Pytorch >= 1.1.0) or [TensorBoardX](https://github.com/lanpa/tensorboardX) (<https://github.com/lanpa/tensorboardX>) - a third party library. Add code in `trainer.py` to visualize the testing MAP and training loss in Tensorboard. *You may have to reload the kernel for these changes to take effect*

Show clear screenshots of the learning curves of testing MAP and training loss for 5 epochs (batch size=20, learning rate=0.001). Please evaluate your model to calculate the MAP on the testing dataset every 100 iterations.

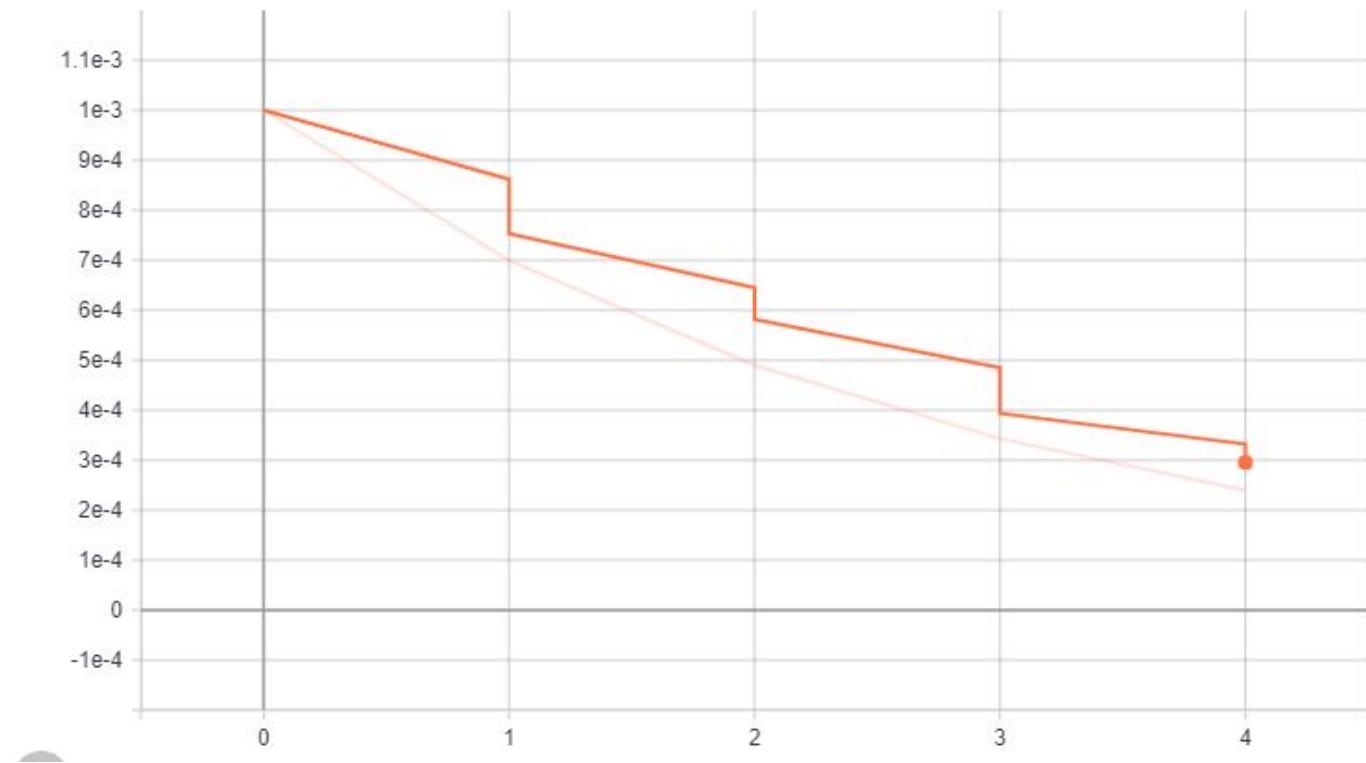
```
In [2]: args = ARGs(epochs=5, batch_size=20, lr=0.001)
model = SimpleCNN(num_classes=len(VOCdataset.CLASS_NAMES), inp_size=64, c_dim=3)
optimizer = torch.optim.Adam(model.parameters(), lr=args.lr)
scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=1, gamma=args.gamma)
if __name__ == "__main__":
    test_ap, test_map = trainer.train(args, model, optimizer, scheduler)
    print('test map:', test_map)
```

```
Train Epoch: 0 [0 (0%)] Loss: 0.694451 | mAP: 0.076841
Train Epoch: 0 [100 (40%)] Loss: 0.240711 | mAP: 0.123666
Train Epoch: 0 [200 (80%)] Loss: 0.242937 | mAP: 0.163159
Train Epoch: 1 [300 (20%)] Loss: 0.212498 | mAP: 0.173409
Train Epoch: 1 [400 (59%)] Loss: 0.211122 | mAP: 0.193135
Train Epoch: 1 [500 (99%)] Loss: 0.216955 | mAP: 0.204963
Train Epoch: 2 [600 (39%)] Loss: 0.186005 | mAP: 0.218065
Train Epoch: 2 [700 (79%)] Loss: 0.140446 | mAP: 0.224668
Train Epoch: 3 [800 (19%)] Loss: 0.208131 | mAP: 0.232051
Train Epoch: 3 [900 (59%)] Loss: 0.194046 | mAP: 0.234823
Train Epoch: 3 [1000 (98%)] Loss: 0.199973 | mAP: 0.238088
Train Epoch: 4 [1100 (38%)] Loss: 0.199175 | mAP: 0.241002
Train Epoch: 4 [1200 (78%)] Loss: 0.167286 | mAP: 0.246116
test map: 0.24888071523104266
```

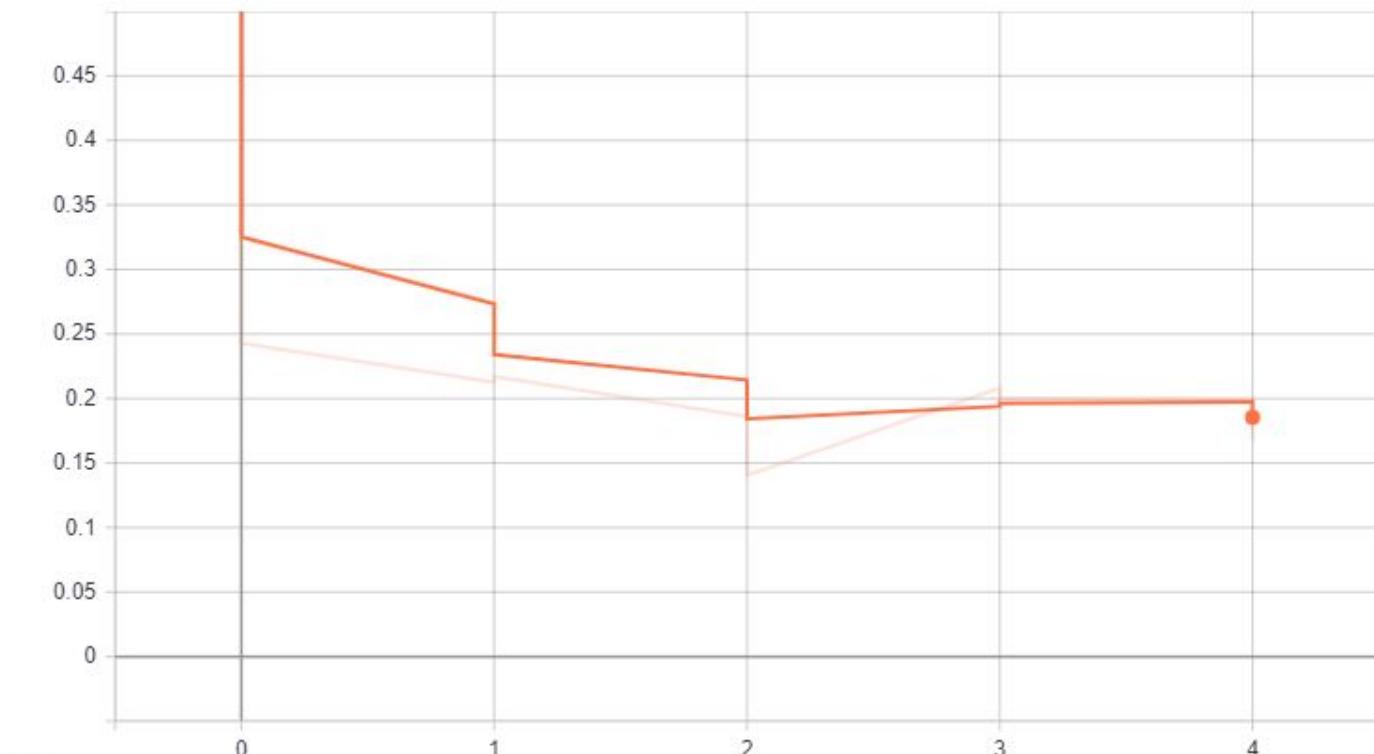
## INSERT YOUR TENSORBOARD SCREENSHOTS HERE

The figures below display the learning rate, training loss, and training map, respectively:

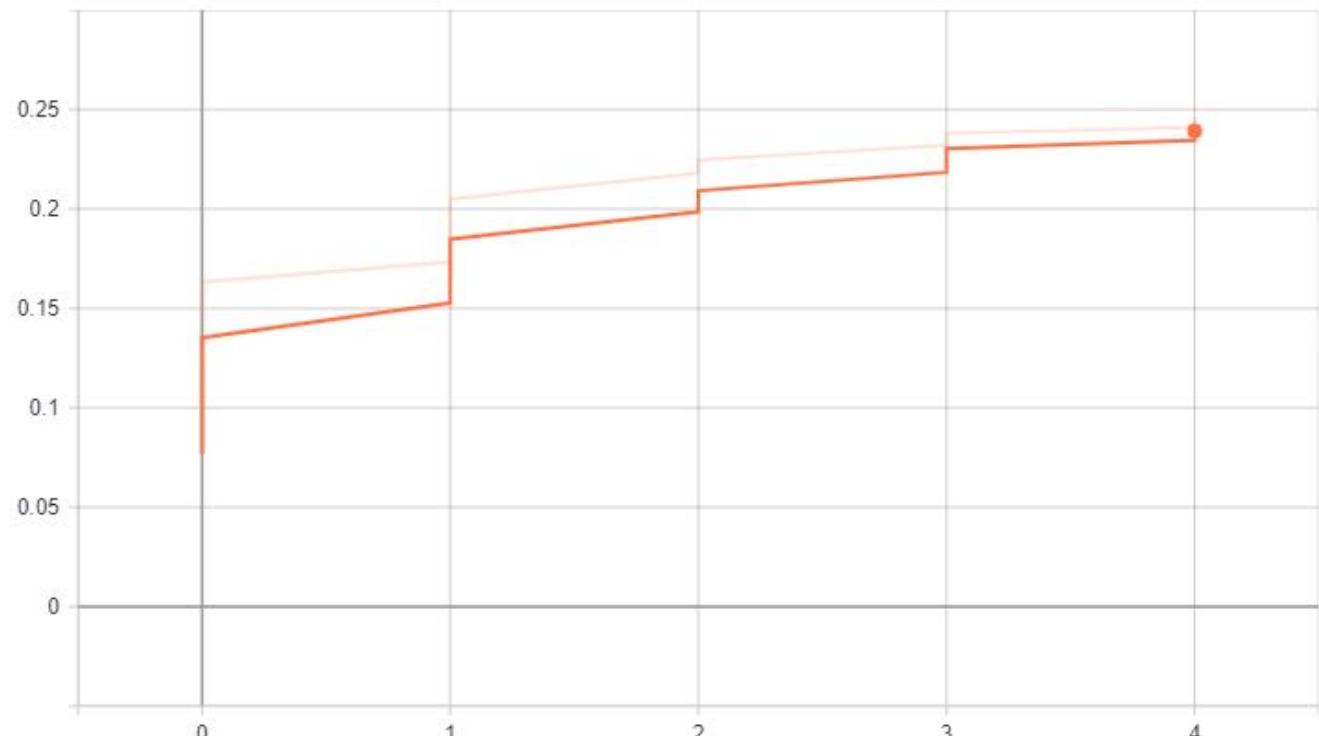
Learning Rate



Loss



mAP



In [ ]:

## Q2: Lets go deeper! CaffeNet for PASCAL classification (20 pts)

**Note:** You are encouraged to reuse code from the previous task. Finish Q1 if you haven't already!

As you might have seen, the performance of the SimpleCNN model was pretty low for PASCAL. This is expected as PASCAL is much more complex than FASHION MNIST, and we need a much beefier model to handle it.

In this task we will be constructing a variant of the [AlexNet \(<https://papers.nips.cc/paper/4824-imagenet-classification-with-deep-convolutional-neural-networks.pdf>\)](https://papers.nips.cc/paper/4824-imagenet-classification-with-deep-convolutional-neural-networks.pdf) architecture, known as CaffeNet. If you are familiar with Caffe, a prototxt of the network is available [here \(\[https://github.com/BVLC/caffe/blob/master/models/bvlc\\\_reference\\\_caffenet/train\\\_val.prototxt\]\(https://github.com/BVLC/caffe/blob/master/models/bvlc\_reference\_caffenet/train\_val.prototxt\)\)](https://github.com/BVLC/caffe/blob/master/models/bvlc_reference_caffenet/train_val.prototxt). A visualization of the network is available [here \(<http://ethereon.github.io/netscope/#/preset/caffenet>\)](http://ethereon.github.io/netscope/#/preset/caffenet).

### 2.1 Build CaffeNet (5 pts)

Here is the exact model we want to build. In this task, `torchvision.models.xxx()` is NOT allowed. Define your own CaffeNet! We use the following operator notation for the architecture:

1. Convolution: A convolution with kernel size  $k$ , stride  $s$ , output channels  $n$ , padding  $p$  is represented as  $\text{conv}(k, s, n, p)$ .
2. Max Pooling: A max pool operation with kernel size  $k$ , stride  $s$  as  $\text{maxpool}(k, s)$ .
3. Fully connected: For  $n$  output units,  $\text{FC}(n)$ .
4. ReLU: For rectified linear non-linearity  $\text{relu}()$

ARCHITECTURE:

```
-> image
-> conv(11, 4, 96, 'VALID')
-> relu()
-> max_pool(3, 2)
-> conv(5, 1, 256, 'SAME')
-> relu()
-> max_pool(3, 2)
-> conv(3, 1, 384, 'SAME')
-> relu()
-> conv(3, 1, 384, 'SAME')
-> relu()
-> conv(3, 1, 256, 'SAME')
-> relu()
-> max_pool(3, 2)
-> flatten()
-> fully_connected(4096)
-> relu()
-> dropout(0.5)
-> fully_connected(4096)
-> relu()
-> dropout(0.5)
-> fully_connected(20)
```

In [ ]:

```
import torch
import torch.nn as nn
import torch.nn.functional as F
import matplotlib.pyplot as plt
# %matplotlib inline

import trainer
from utils import ARGS
from simple_cnn import SimpleCNN
from voc_dataset import VOCDataset

def get_fc(inp_dim, out_dim, non_linear='relu'):
    """
    Mid-Level API. It is useful to customize your own for Large code repo.
    :param inp_dim: int, intput dimension
    :param out_dim: int, output dimension
    :param non_linear: str, 'relu', 'softmax'
    :return: list of layers [FC(inp_dim, out_dim), (non Linear Layer)]
    """
    layers = []
    layers.append(nn.Linear(inp_dim, out_dim))
    if non_linear == 'relu':
        layers.append(nn.ReLU())
    elif non_linear == 'softmax':
        layers.append(nn.Softmax(dim=1))
    elif non_linear == 'none':
        pass
    else:
        raise NotImplementedError
    return layers

class CaffeNet(nn.Module):
    def __init__(self):
        super().__init__()
        c_dim = 3
        self.conv1 = nn.Conv2d(c_dim, 96, 11, 4, padding=0) # valid padding
        self.pool1 = nn.MaxPool2d(3, 2)
        self.conv2 = nn.Conv2d(96, 256, 5, padding=2) # same padding
        self.pool2 = nn.MaxPool2d(3, 2)
        self.conv3 = nn.Conv2d(256, 384, 3, padding=1) # same padding
        self.conv4 = nn.Conv2d(384, 384, 3, padding=1) # same padding
        self.conv5 = nn.Conv2d(384, 256, 3, padding=1) # same padding
        self.pool3 = nn.MaxPool2d(3, 2)
        self.flat_dim = 5*5*256 # replace with the actual value
        self.fc1 = nn.Sequential(*get_fc(self.flat_dim, 4096, 'relu'))
        self.dropout1 = nn.Dropout(p=0.5)
        self.fc2 = nn.Sequential(*get_fc(4096, 4096, 'relu'))
        self.dropout2 = nn.Dropout(p=0.5)
        self.fc3 = nn.Sequential(*get_fc(4096, 20, 'none'))

        self.nonlinear = lambda x: torch.clamp(x, 0)

    def forward(self, x):
        N = x.size(0)
        x = self.conv1(x)
```

```

        x = self.nonlinear(x)
        x = self.pool1(x)

        x = self.conv2(x)
        x = self.nonlinear(x)
        x = self.pool2(x)

        x = self.conv3(x)
        x = self.nonlinear(x)
        x = self.conv4(x)
        x = self.nonlinear(x)
        x = self.conv5(x)
        x = self.nonlinear(x)
        x = self.pool3(x)
        x = x.view(N, self.flat_dim) # flatten the array

        out = self.fc1(x)
        out = self.nonlinear(out)
        out = self.dropout1(out)
        out = self.fc2(out)
        out = self.nonlinear(out)
        out = self.dropout2(out)
        out = self.fc3(out)

    return out

```

## 2.2 Save the Model (5 pts)

Finish code stubs for saving the model periodically into `trainer.py`. You will need these models later

## 2.3 Train and Test (5pts)

Show clear screenshots of testing MAP and training loss for 50 epochs. Please evaluate your model to calculate the MAP on the testing dataset every 250 iterations. Use the following hyperparamters:

- batch\_size=32
- Adam optimizer with lr=0.0001

**NOTE: SAVE AT LEAST 5 EVENLY SPACED CHECKPOINTS DURING TRAINING (1 at end)**

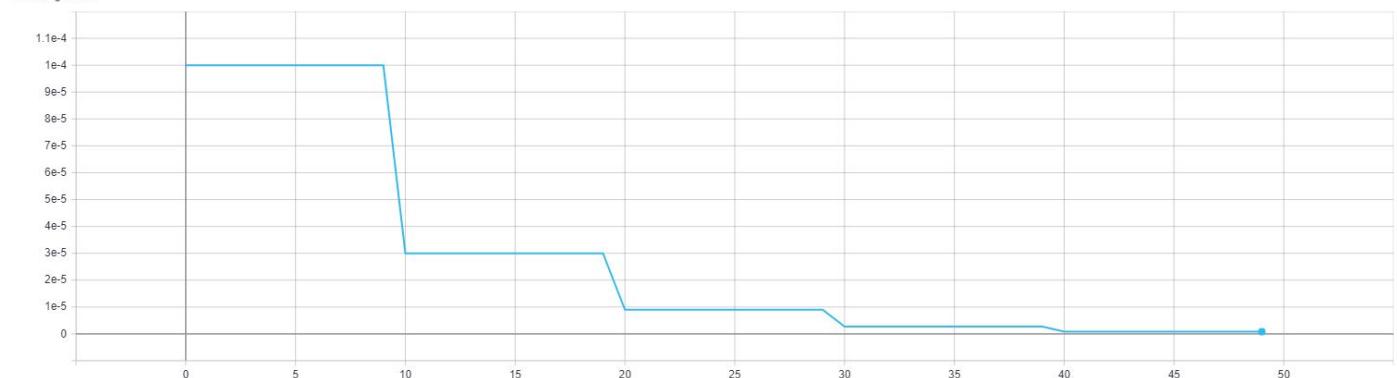
```
In [ ]: def xavier_normal_init(m):
    if(type(m)==nn.Conv1d or type(m)==nn.Conv2d or type(m)==nn.Linear):
        torch.nn.init.xavier_normal_(m.weight.data)
        if(m.bias is not None):
            torch.nn.init.xavier_normal_(m.weight.data)
```

```
In [ ]: args = ARGS(batch_size = 32, epochs=50, lr = 0.0001)
args.gamma = 0.3
weightDecay = 5e-5
model = CaffeNet()
model.apply(xavier_normal_init)
optimizer = torch.optim.Adam(model.parameters(), lr = args.lr, weight_decay=weightDecay) # ,weight_decay=weightDecay
scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=10, gamma=args.gamma)
if __name__ == '__main__':
    test_ap, test_map = trainer.train(args, model, optimizer, scheduler)
    print('test map:', test_map)
```

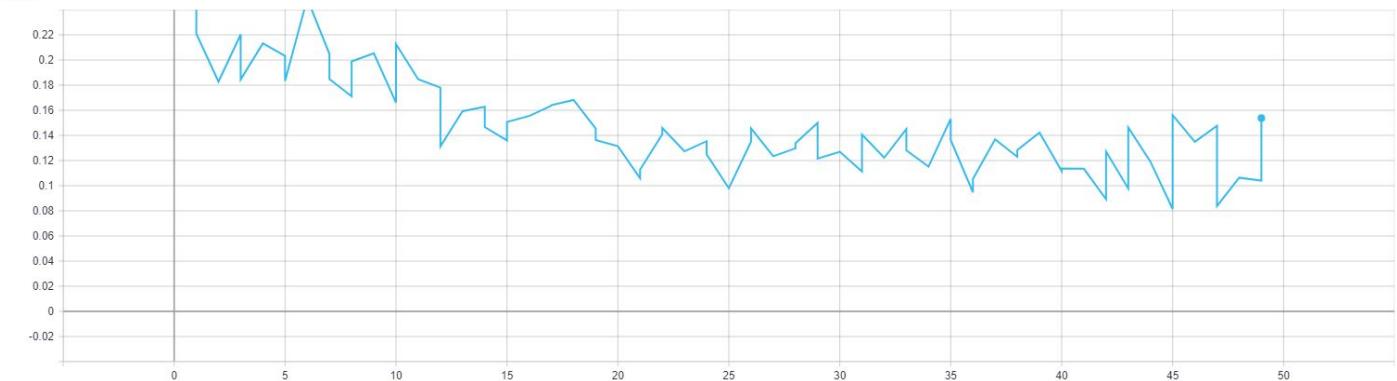
## INSERT YOUR TENSORBOARD SCREENSHOTS HERE

The figures below display the learning rate, loss, and map, respectively, during training.

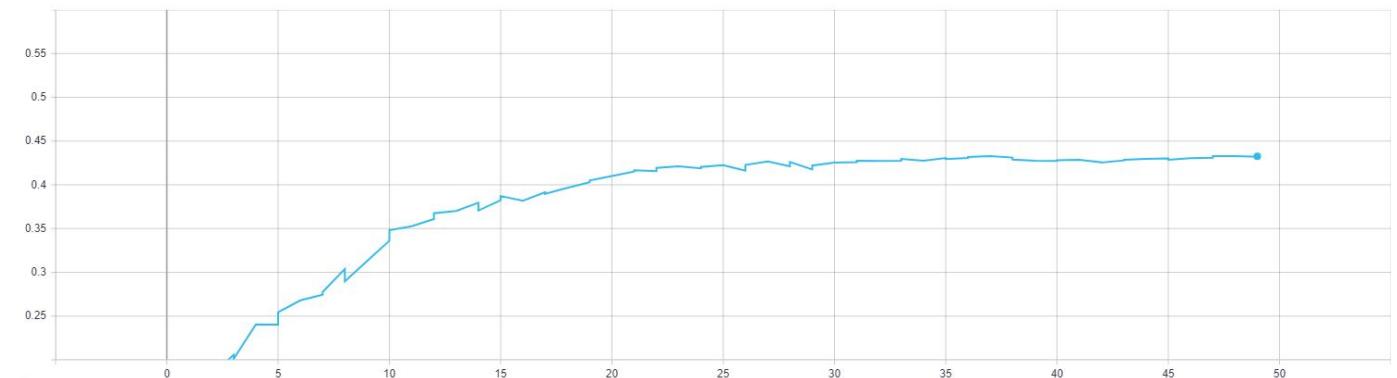
Learning\_Rate



Loss



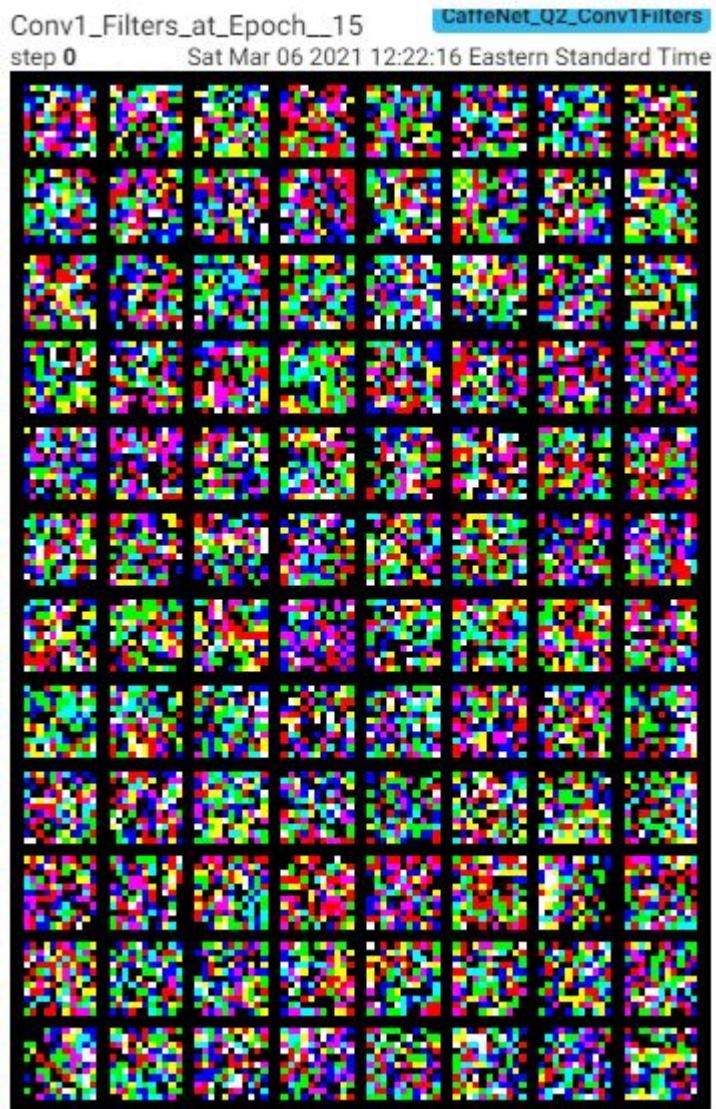
mAP



## **2.4 Visualizing: Conv-1 filters (5pts)**

Extract and compare the conv1 filters, at different stages of the training (at least from 3 different iterations). Show at least 5 filters.

The following two image grids below show the Conv1 layer filters at epochs 15, 30, and 45, respectively:

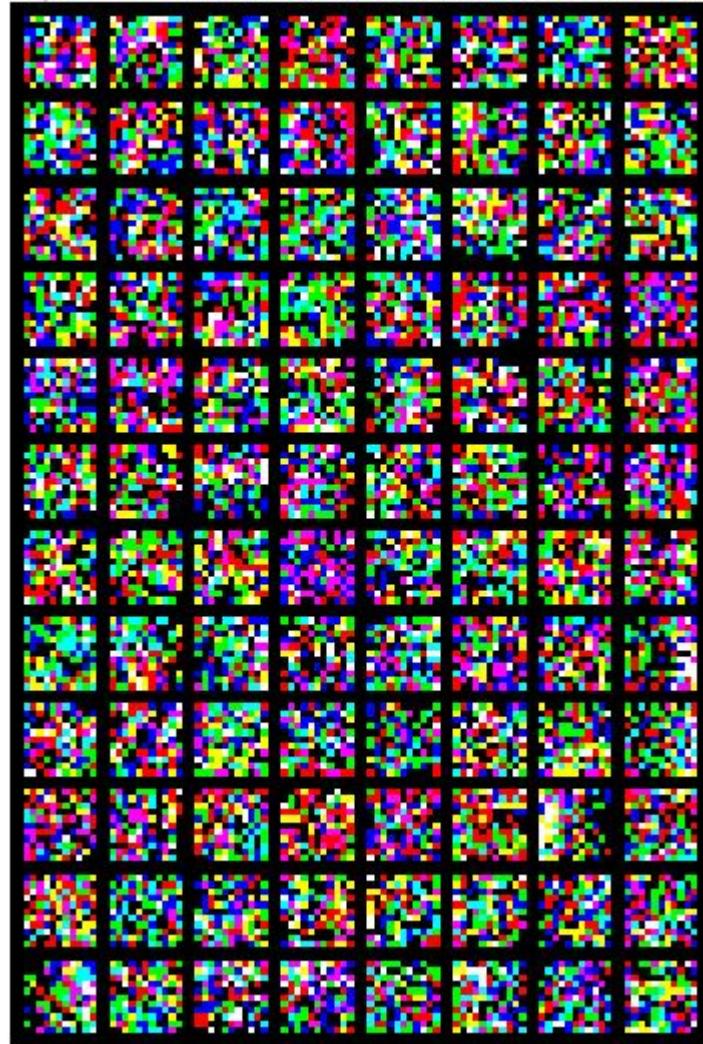


Conv1\_Filters\_at\_Epoch\_30

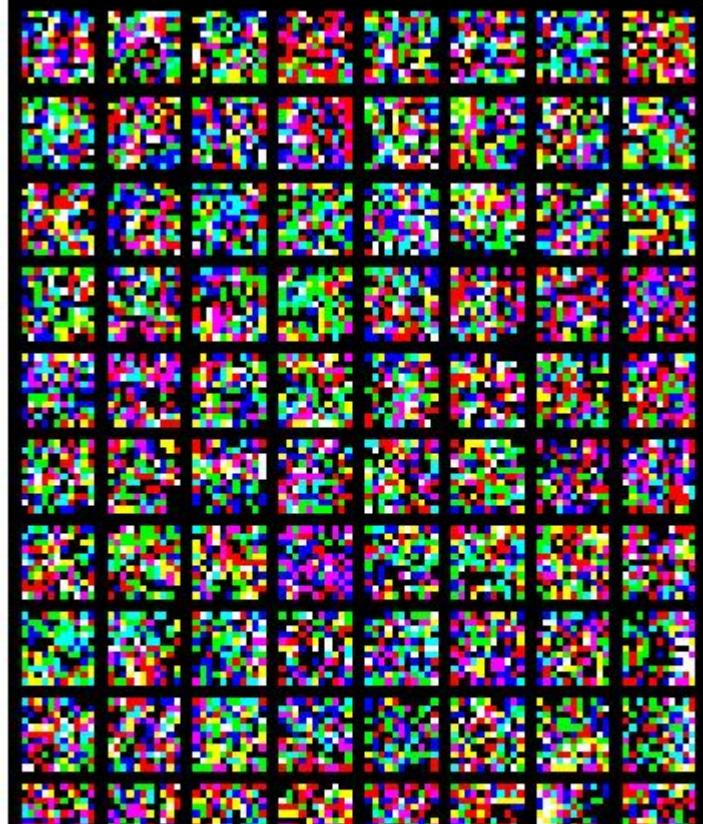
step 0

Sat Mar 06 2021 12:38:59 Eastern Standard Time

CaffeNet\_Q2\_Conv1Filters



Conv1\_Filters\_at\_Epoch\_45      **CaffeNet\_Q2\_Conv1Filters**  
step 0      Sat Mar 06 2021 12:55:07 Eastern Standard Time



## Appendix A: Training Epoch Information

The following log shown below displays the batch iteration, loss, and mAP calculation at various points during the training process:

```
Train Epoch: 0 [0 (0%)] Loss: 0.694260 | mAP: 0.079919
Train Epoch: 0 [100 (64%)] Loss: 0.252534 | mAP: 0.106128
Train Epoch: 1 [200 (27%)] Loss: 0.245160 | mAP: 0.125702
Train Epoch: 1 [300 (91%)] Loss: 0.220970 | mAP: 0.154655
Train Epoch: 2 [400 (55%)] Loss: 0.182752 | mAP: 0.181409
Train Epoch: 3 [500 (18%)] Loss: 0.220406 | mAP: 0.205849
Train Epoch: 3 [600 (82%)] Loss: 0.184534 | mAP: 0.200986
Train Epoch: 4 [700 (46%)] Loss: 0.213212 | mAP: 0.240417
Train Epoch: 5 [800 (10%)] Loss: 0.203305 | mAP: 0.240243
Train Epoch: 5 [900 (73%)] Loss: 0.183403 | mAP: 0.254320
Train Epoch: 6 [1000 (37%)] Loss: 0.248874 | mAP: 0.267995
Train Epoch: 7 [1100 (1%)] Loss: 0.204796 | mAP: 0.274374
Train Epoch: 7 [1200 (64%)] Loss: 0.184777 | mAP: 0.277550
Train Epoch: 8 [1300 (28%)] Loss: 0.171141 | mAP: 0.303782
Train Epoch: 8 [1400 (92%)] Loss: 0.198896 | mAP: 0.289614
Train Epoch: 9 [1500 (55%)] Loss: 0.205384 | mAP: 0.312908
Train Epoch: 10 [1600 (19%)] Loss: 0.166079 | mAP: 0.336100
Train Epoch: 10 [1700 (83%)] Loss: 0.212687 | mAP: 0.348000
Train Epoch: 11 [1800 (46%)] Loss: 0.184672 | mAP: 0.352580
Train Epoch: 12 [1900 (10%)] Loss: 0.178017 | mAP: 0.360864
Train Epoch: 12 [2000 (74%)] Loss: 0.131287 | mAP: 0.367583
Train Epoch: 13 [2100 (38%)] Loss: 0.159262 | mAP: 0.370199
Train Epoch: 14 [2200 (1%)] Loss: 0.162815 | mAP: 0.379512
Train Epoch: 14 [2300 (65%)] Loss: 0.146647 | mAP: 0.370575
Train Epoch: 15 [2400 (29%)] Loss: 0.136011 | mAP: 0.382428
Train Epoch: 15 [2500 (92%)] Loss: 0.150734 | mAP: 0.387179
Train Epoch: 16 [2600 (56%)] Loss: 0.155456 | mAP: 0.381900
Train Epoch: 17 [2700 (20%)] Loss: 0.163627 | mAP: 0.391537
Train Epoch: 17 [2800 (83%)] Loss: 0.164096 | mAP: 0.389797
Train Epoch: 18 [2900 (47%)] Loss: 0.168196 | mAP: 0.396490
Train Epoch: 19 [3000 (11%)] Loss: 0.145444 | mAP: 0.402965
Train Epoch: 19 [3100 (75%)] Loss: 0.136387 | mAP: 0.404979
Train Epoch: 20 [3200 (38%)] Loss: 0.131393 | mAP: 0.410043
Train Epoch: 21 [3300 (2%)] Loss: 0.105851 | mAP: 0.415120
Train Epoch: 21 [3400 (66%)] Loss: 0.112801 | mAP: 0.416542
Train Epoch: 22 [3500 (29%)] Loss: 0.140922 | mAP: 0.415565
Train Epoch: 22 [3600 (93%)] Loss: 0.145867 | mAP: 0.419292
Train Epoch: 23 [3700 (57%)] Loss: 0.127371 | mAP: 0.421104
Train Epoch: 24 [3800 (20%)] Loss: 0.135417 | mAP: 0.418907
Train Epoch: 24 [3900 (84%)] Loss: 0.124882 | mAP: 0.420355
Train Epoch: 25 [4000 (48%)] Loss: 0.097941 | mAP: 0.422405
Train Epoch: 26 [4100 (11%)] Loss: 0.134959 | mAP: 0.416246
Train Epoch: 26 [4200 (75%)] Loss: 0.145727 | mAP: 0.422733
```

Train Epoch: 27 [4300 (39%)] Loss: 0.123460 | mAP: 0.426791  
Train Epoch: 28 [4400 (3%)] Loss: 0.129733 | mAP: 0.421078  
Train Epoch: 28 [4500 (66%)] Loss: 0.133712 | mAP: 0.426055  
Train Epoch: 29 [4600 (30%)] Loss: 0.149955 | mAP: 0.417618  
Train Epoch: 29 [4700 (94%)] Loss: 0.121479 | mAP: 0.421980  
Train Epoch: 30 [4800 (57%)] Loss: 0.127036 | mAP: 0.425303  
Train Epoch: 31 [4900 (21%)] Loss: 0.111305 | mAP: 0.425996  
Train Epoch: 31 [5000 (85%)] Loss: 0.140682 | mAP: 0.427547  
Train Epoch: 32 [5100 (48%)] Loss: 0.122193 | mAP: 0.427262  
Train Epoch: 33 [5200 (12%)] Loss: 0.145076 | mAP: 0.427588  
Train Epoch: 33 [5300 (76%)] Loss: 0.128083 | mAP: 0.429591  
Train Epoch: 34 [5400 (39%)] Loss: 0.115151 | mAP: 0.427573  
Train Epoch: 35 [5500 (3%)] Loss: 0.153220 | mAP: 0.430578  
Train Epoch: 35 [5600 (67%)] Loss: 0.136315 | mAP: 0.429386  
Train Epoch: 36 [5700 (31%)] Loss: 0.094817 | mAP: 0.430634  
Train Epoch: 36 [5800 (94%)] Loss: 0.105419 | mAP: 0.431691  
Train Epoch: 37 [5900 (58%)] Loss: 0.136911 | mAP: 0.432678  
Train Epoch: 38 [6000 (22%)] Loss: 0.122926 | mAP: 0.431091  
Train Epoch: 38 [6100 (85%)] Loss: 0.128397 | mAP: 0.428878  
Train Epoch: 39 [6200 (49%)] Loss: 0.142264 | mAP: 0.427402  
Train Epoch: 40 [6300 (13%)] Loss: 0.111605 | mAP: 0.427158  
Train Epoch: 40 [6400 (76%)] Loss: 0.113659 | mAP: 0.428018  
Train Epoch: 41 [6500 (40%)] Loss: 0.113470 | mAP: 0.428675  
Train Epoch: 42 [6600 (4%)] Loss: 0.089273 | mAP: 0.425716  
Train Epoch: 42 [6700 (68%)] Loss: 0.127156 | mAP: 0.425413  
Train Epoch: 43 [6800 (31%)] Loss: 0.097778 | mAP: 0.427899  
Train Epoch: 43 [6900 (95%)] Loss: 0.146353 | mAP: 0.428698  
Train Epoch: 44 [7000 (59%)] Loss: 0.119124 | mAP: 0.429632  
Train Epoch: 45 [7100 (22%)] Loss: 0.081511 | mAP: 0.430262  
Train Epoch: 45 [7200 (86%)] Loss: 0.156132 | mAP: 0.428628  
Train Epoch: 46 [7300 (50%)] Loss: 0.134881 | mAP: 0.430437  
Train Epoch: 47 [7400 (13%)] Loss: 0.147624 | mAP: 0.430969  
Train Epoch: 47 [7500 (77%)] Loss: 0.083810 | mAP: 0.432749  
Train Epoch: 48 [7600 (41%)] Loss: 0.106350 | mAP: 0.432845  
Train Epoch: 49 [7700 (4%)] Loss: 0.104012 | mAP: 0.431998  
Train Epoch: 49 [7800 (68%)] Loss: 0.153784 | mAP: 0.432565  
test map: 0.41263213323264036

## **Q3: Even deeper! Resnet18 for PASCAL classification (15 pts)**

Hopefully we all got much better accuracy with the deeper model! Since 2012, much deeper architectures have been proposed. [ResNet](https://arxiv.org/abs/1512.03385) (<https://arxiv.org/abs/1512.03385>) is one of the popular ones. In this task, we attempt to further improve the performance with the “very deep” ResNet-18 architecture.

### **3.1 Build ResNet-18 (1 pts)**

Write a network modules for the Resnet-18 architecture (refer to the original paper). You can use `torchvision.models` for this section, so it should be very easy!

```
In [ ]: import torch
import torch.nn as nn
import torch.nn.functional as F
from torchvision import models
import matplotlib.pyplot as plt
%matplotlib inline

import trainer
from utils import ARGS
from simple_cnn import SimpleCNN
from voc_dataset import VOCDataset
```

## 3.2 Add Tensorboard Summaries (6 pts)

You should've already written tensorboard summary generation code into `trainer.py` from q1. However, you probably just added the most basic summary features. Please implement the more advanced summaries listed here:

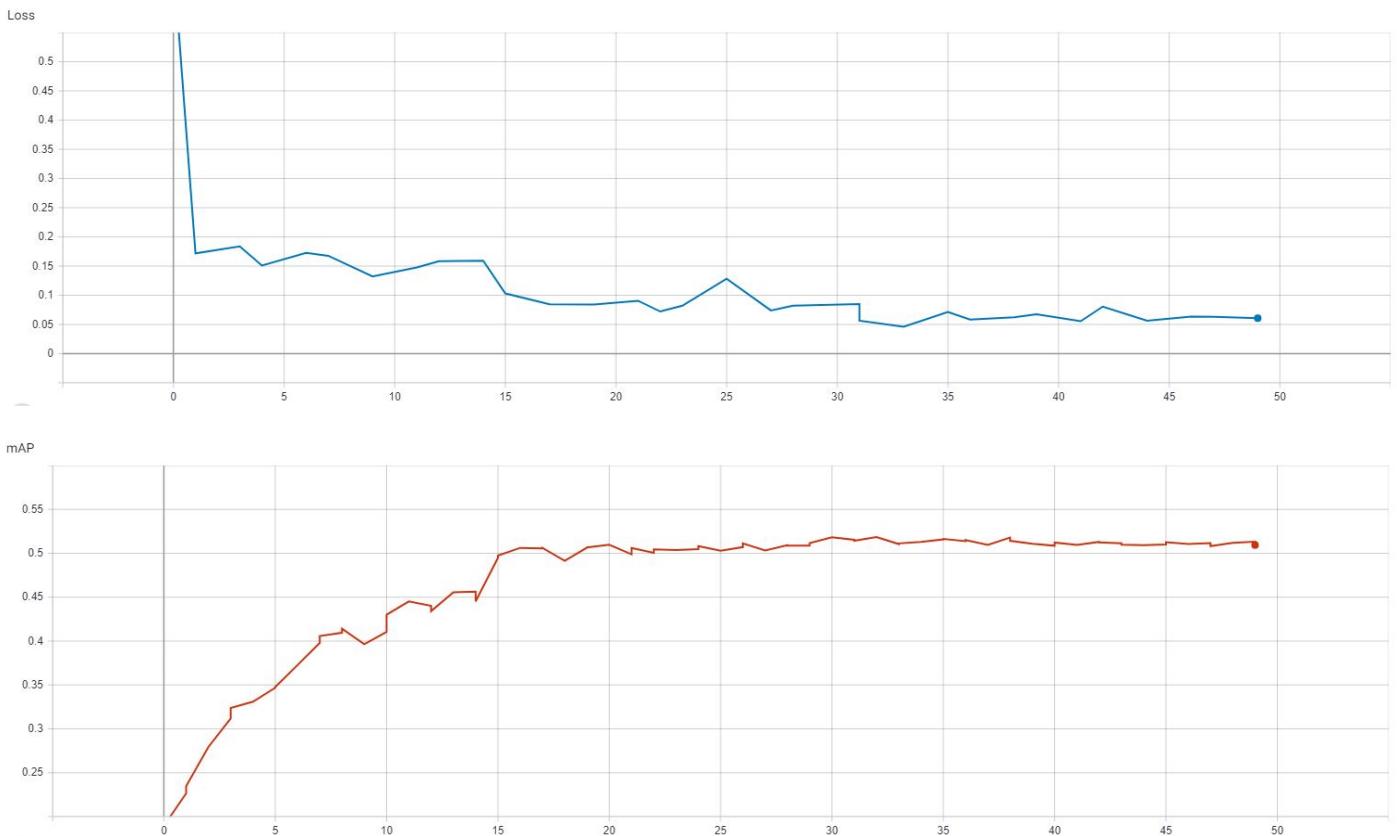
- training loss (should be done)
- testing MAP curves (should be done)
- learning rate
- histogram of gradients

## 3.3 Train and Test (8 pts)

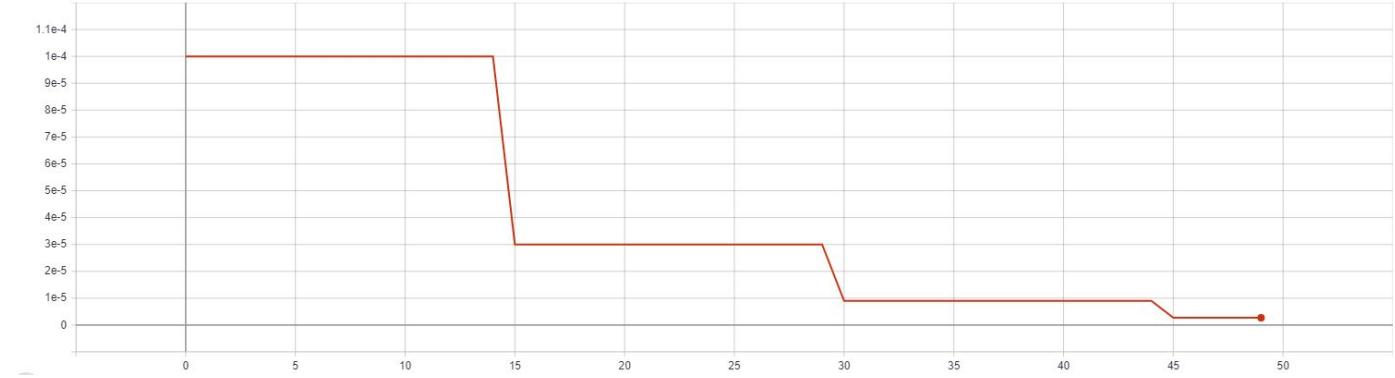
Use the same hyperparameter settings from Task 2, and train the model for 50 epochs. Report tensorboard screenshots for *all* of the summaries listed above (for image summaries show screenshots at  $n \geq 3$  iterations)

**REMEMBER TO SAVE A MODEL AT THE END OF TRAINING**

The figure shown below display the training loss, training mAP, learning rate, and histogram of gradients as displayed through TensorBoard.



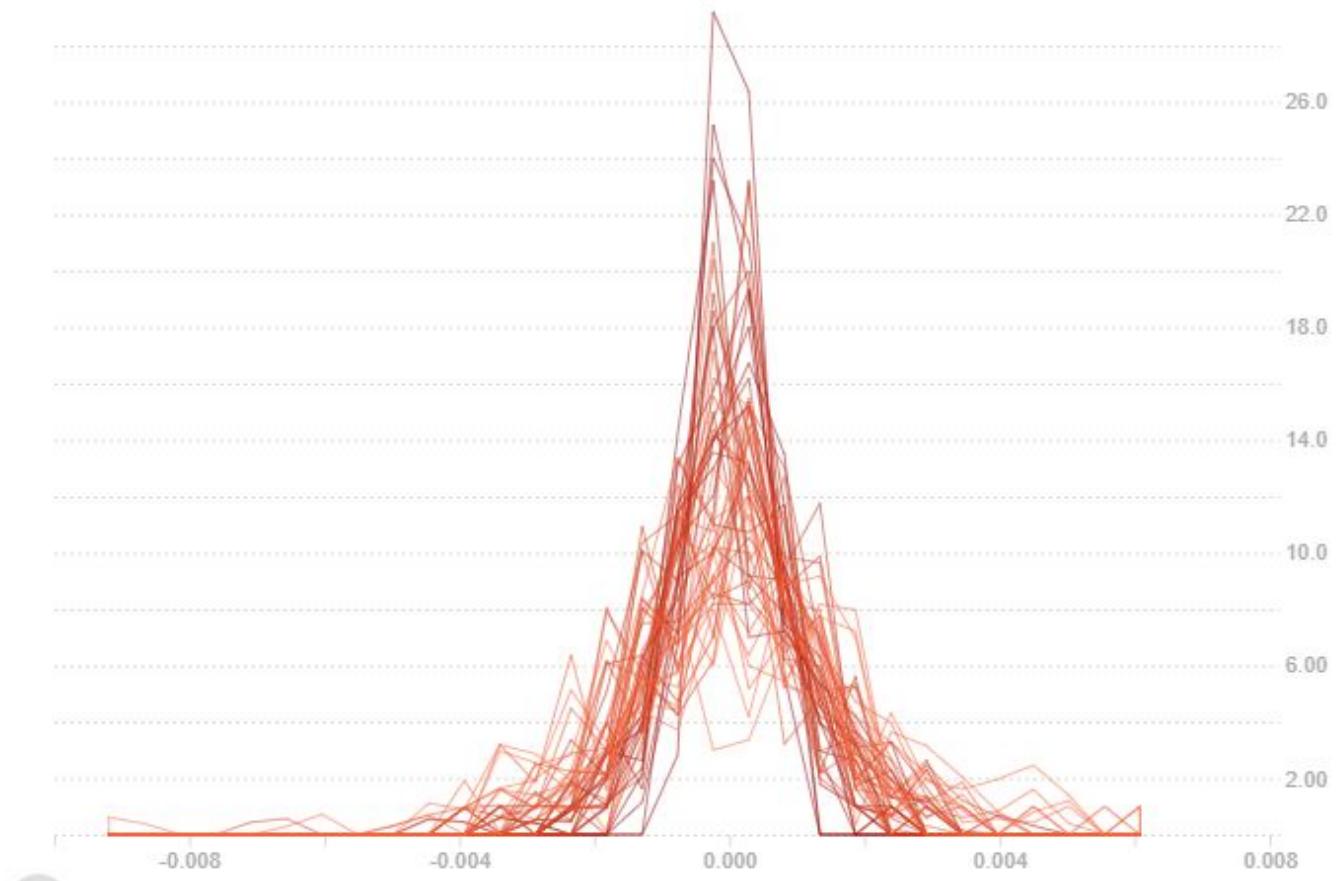
Learning\_Rate



The following 4 images below showcase histogram of gradients for the first two layers:

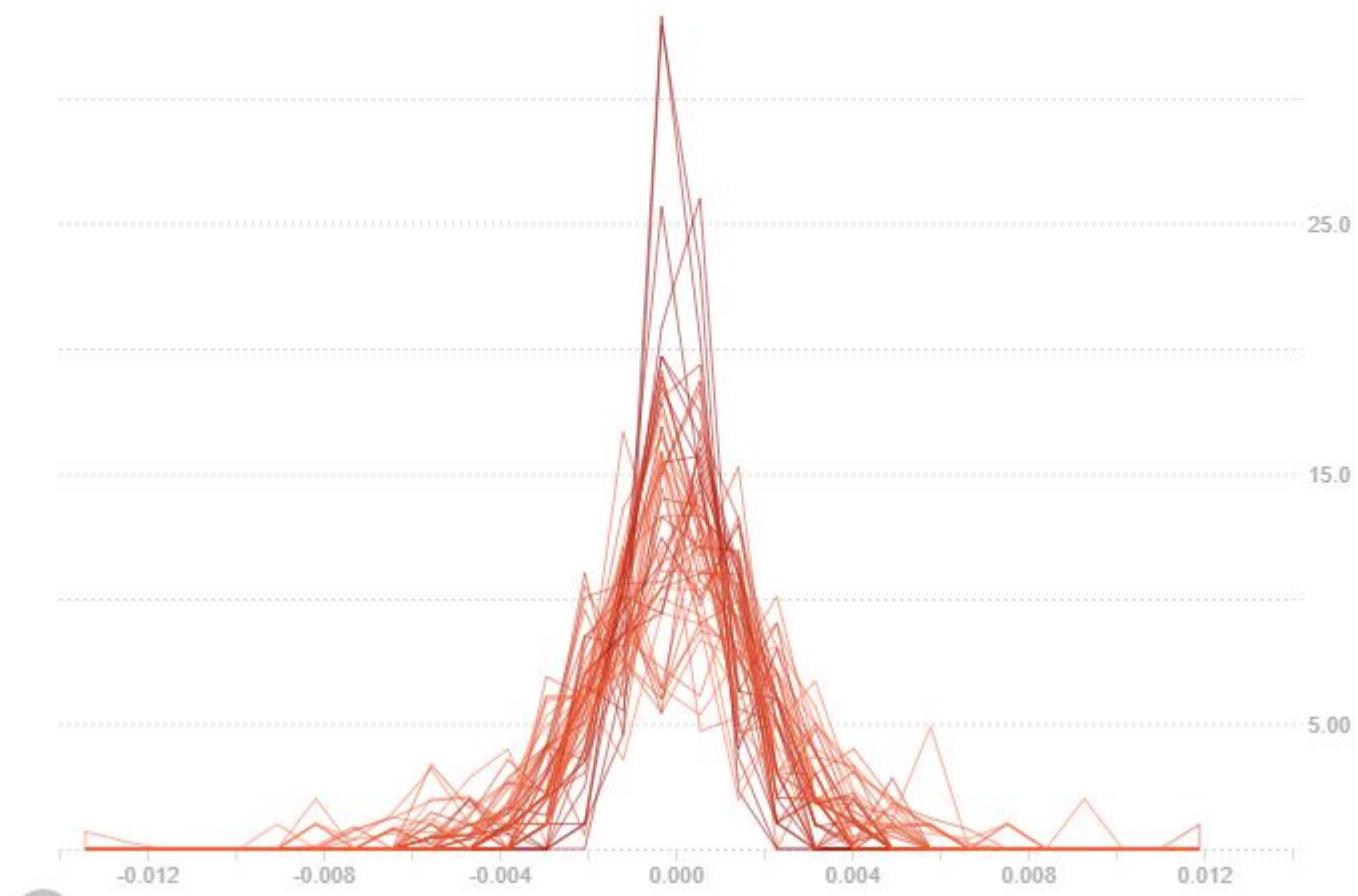
0.bn1.bias.grad

ResNet\_Q3\_histogram



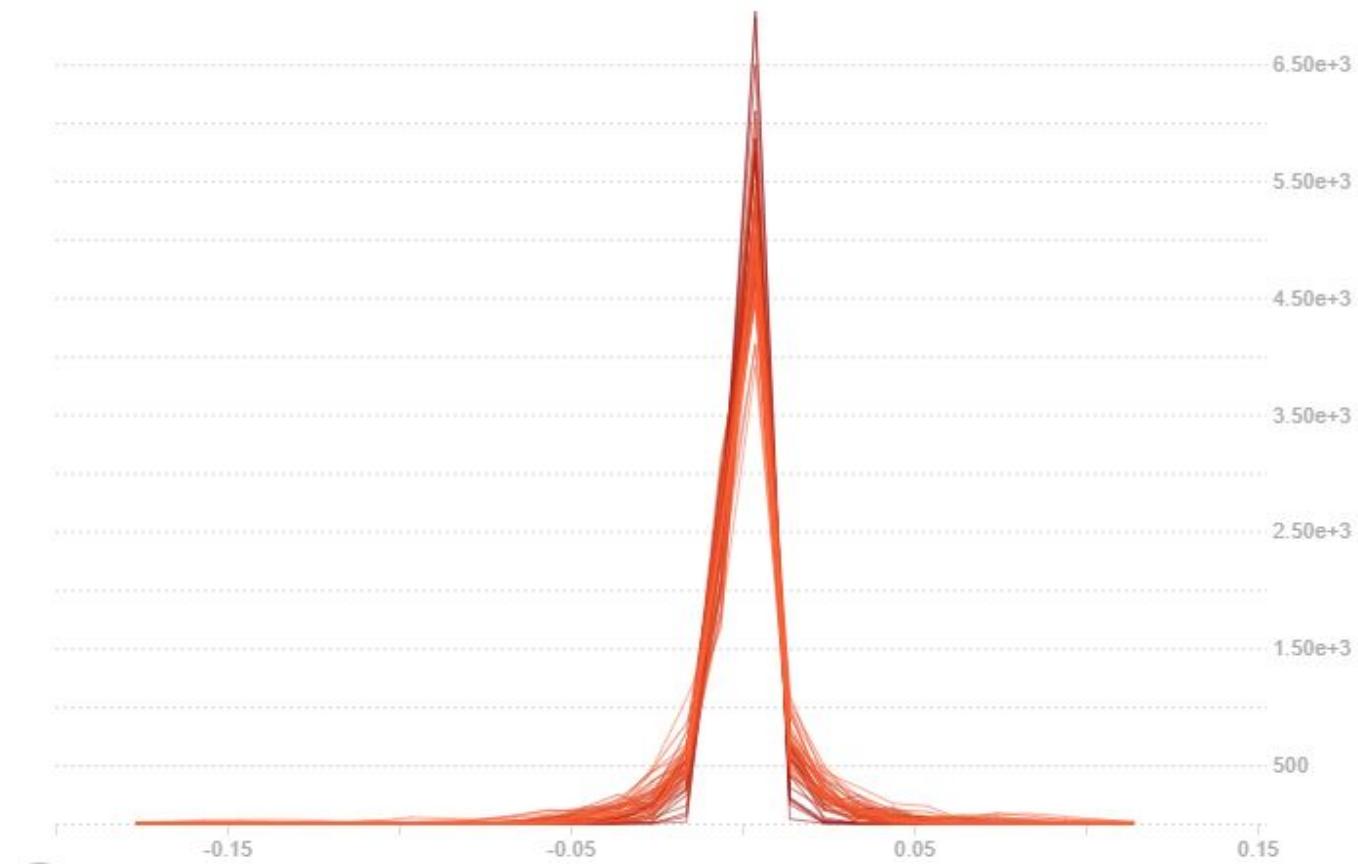
0.bn1.weight.grad

ResNet\_Q3\_histogram



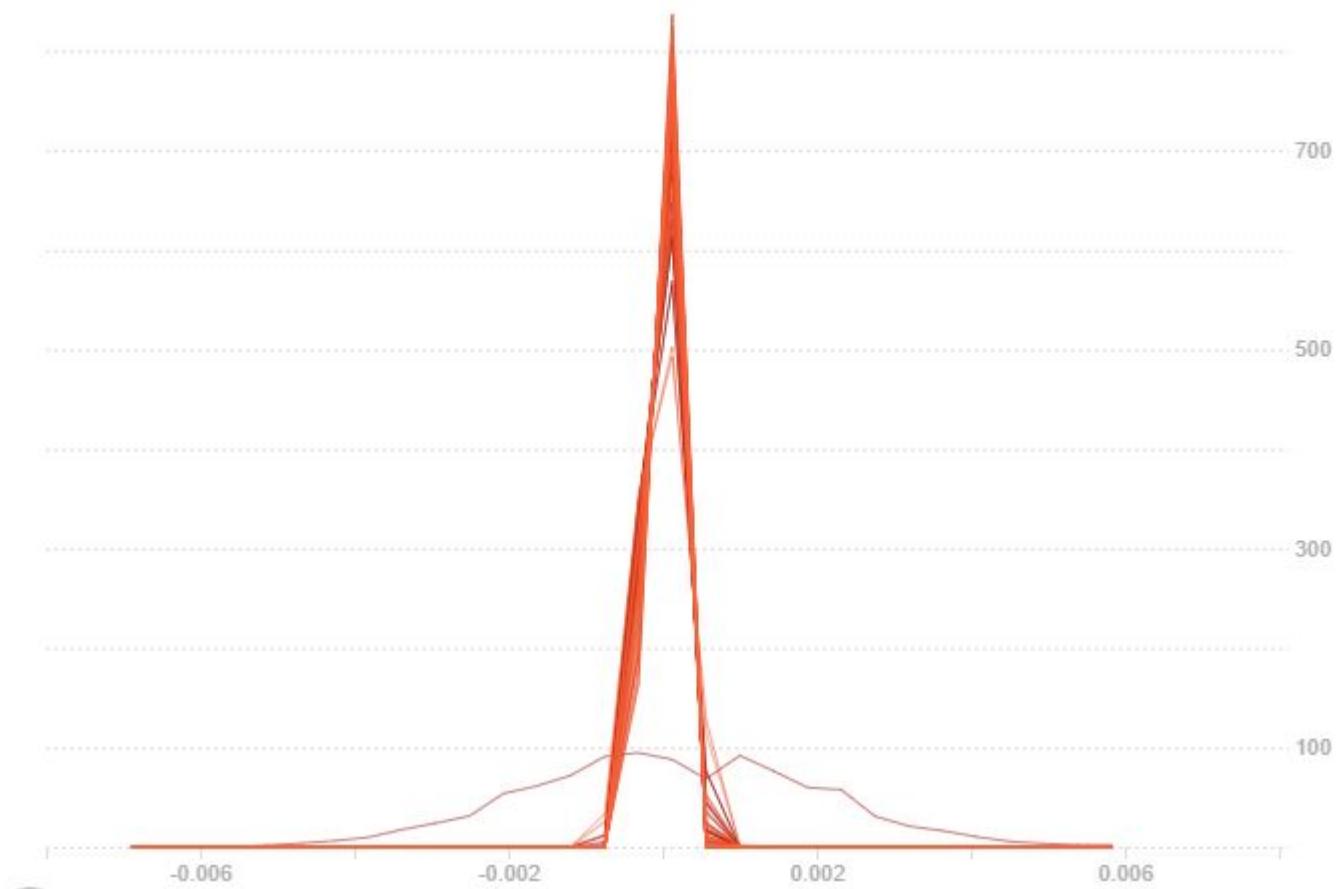
0.conv1.weight.grad

ResNet\_Q3\_histogram



0.fc.bias.grad

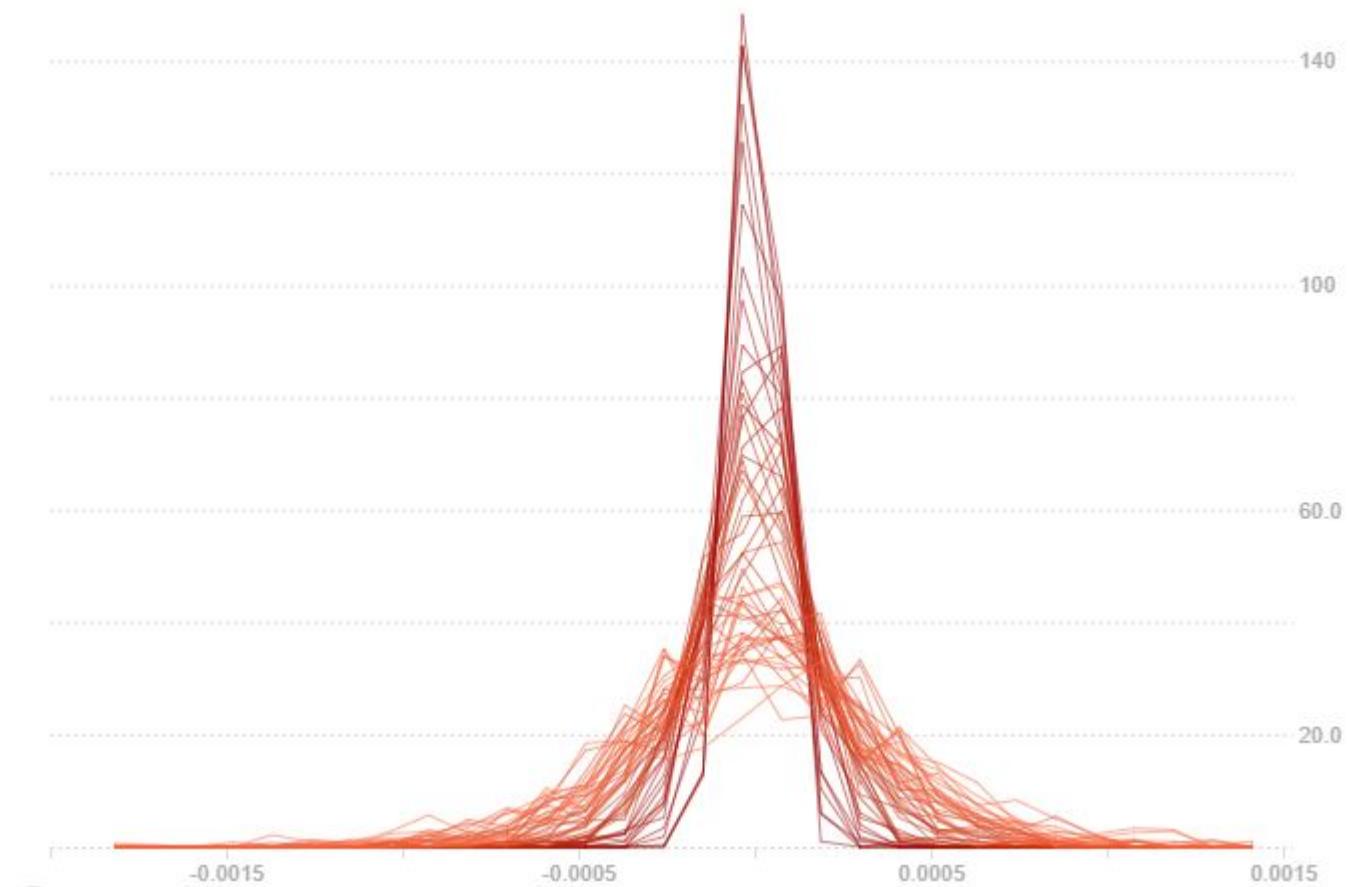
ResNet\_Q3\_histogram



The following 4 images below showcase histogram of gradients for the middle two layers:

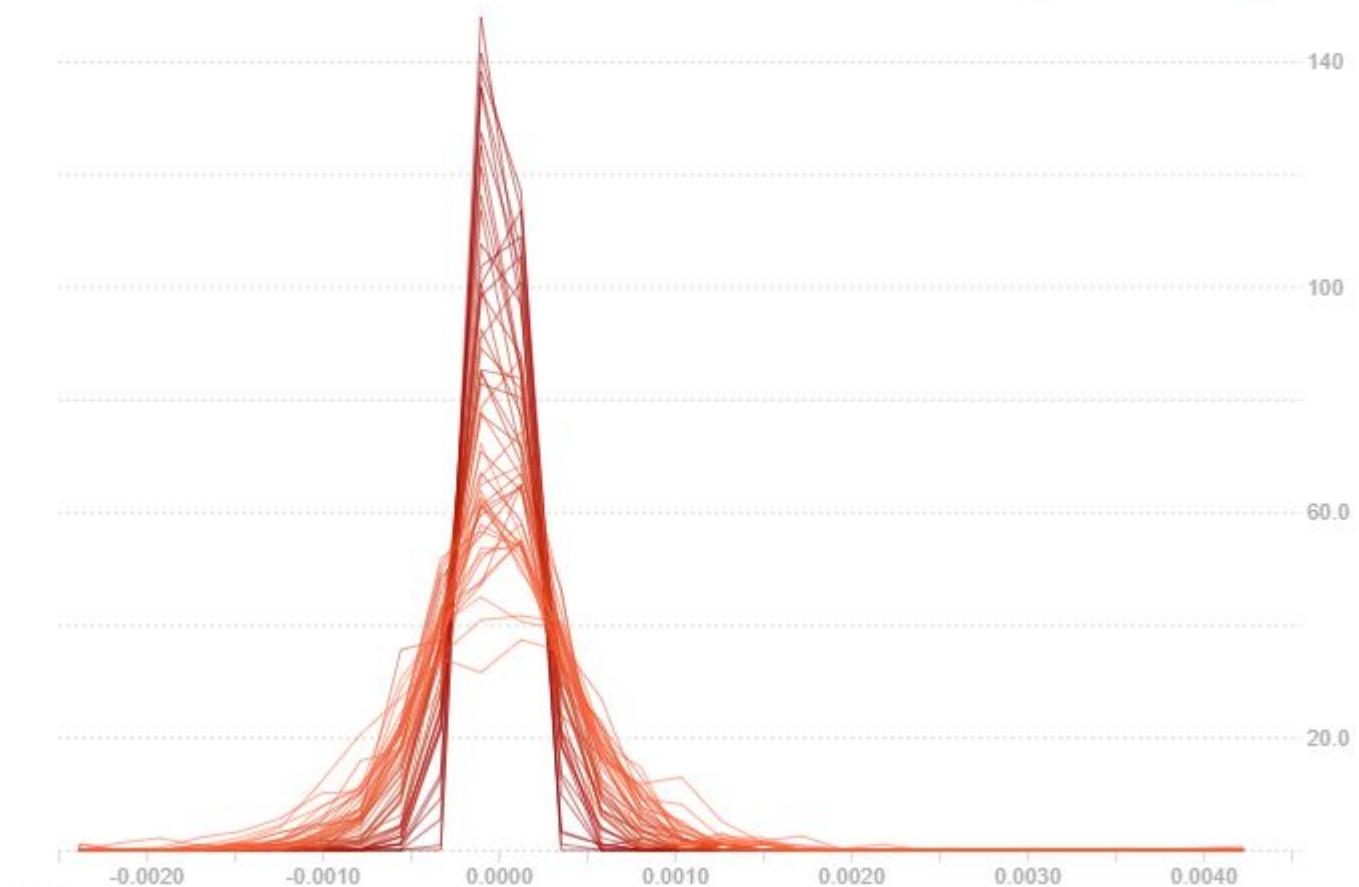
0.layer3.0.bn2.bias.grad

ResNet\_Q3\_histogram



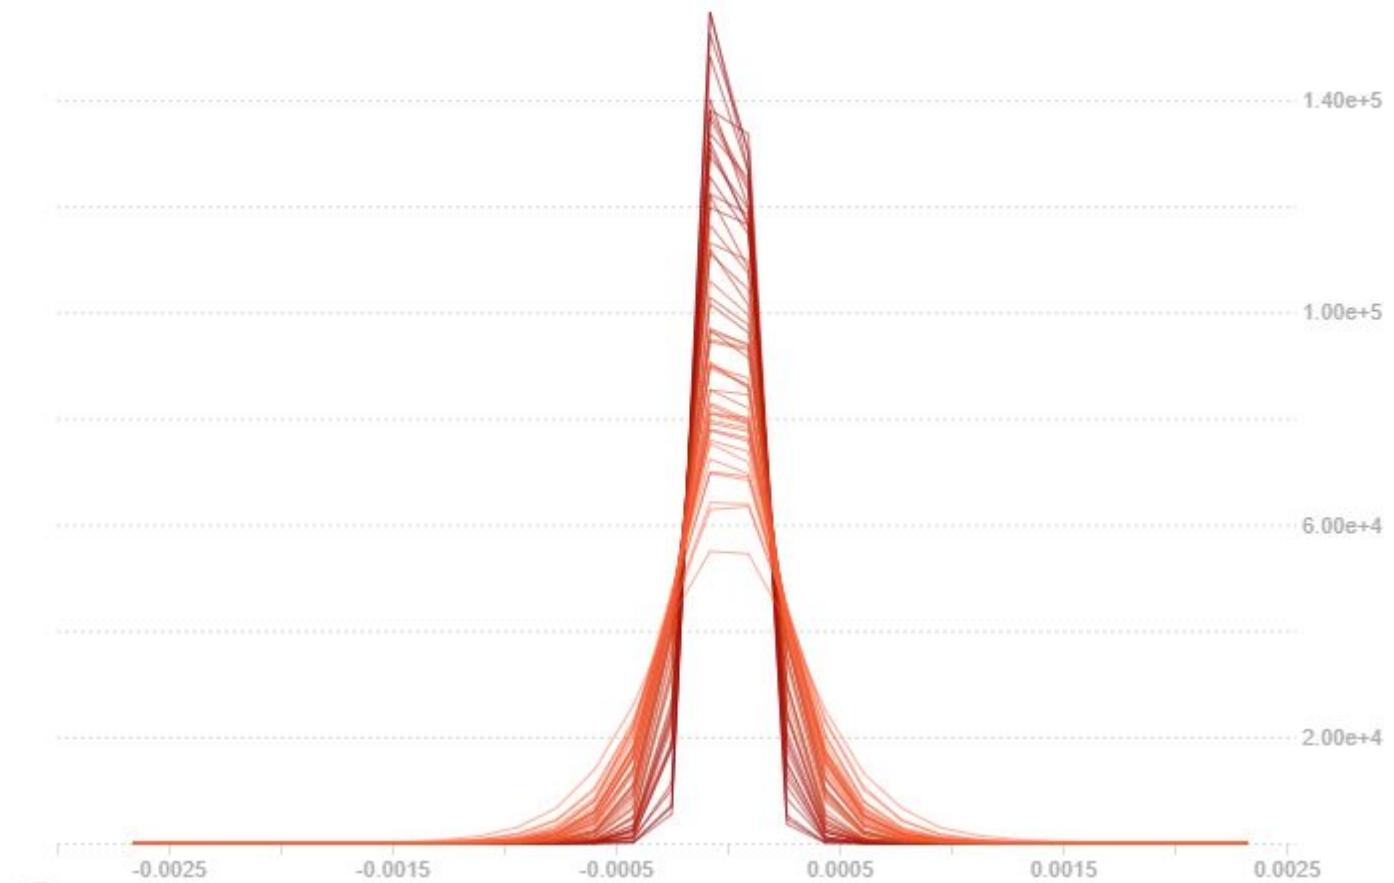
0.layer3.0.bn2.weight.grad

ResNet\_Q3\_histogram



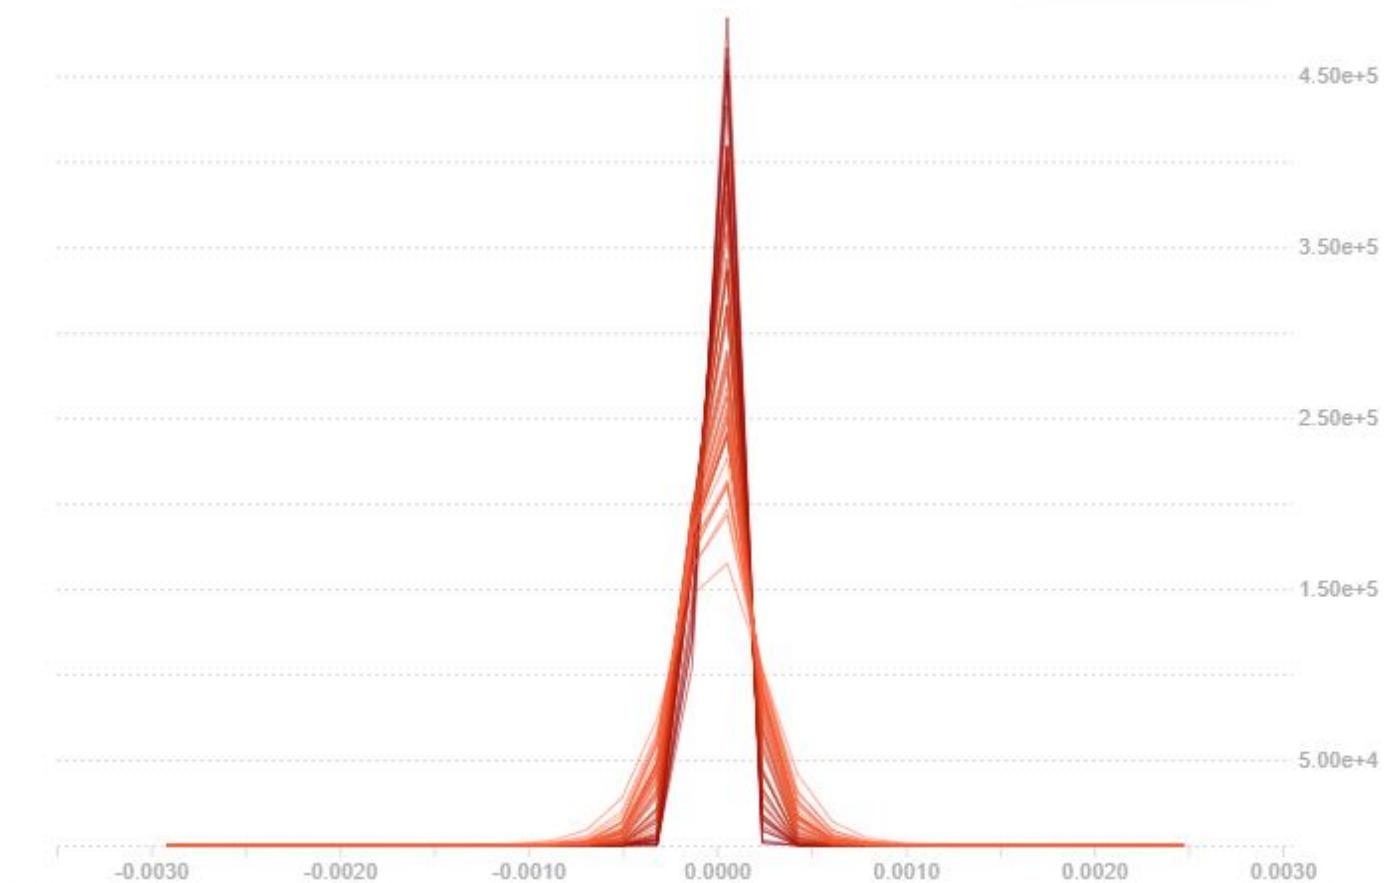
0.layer3.0.conv1.weight.grad

ResNet\_Q3\_histogram



0.layer3.0.conv2.weight.grad

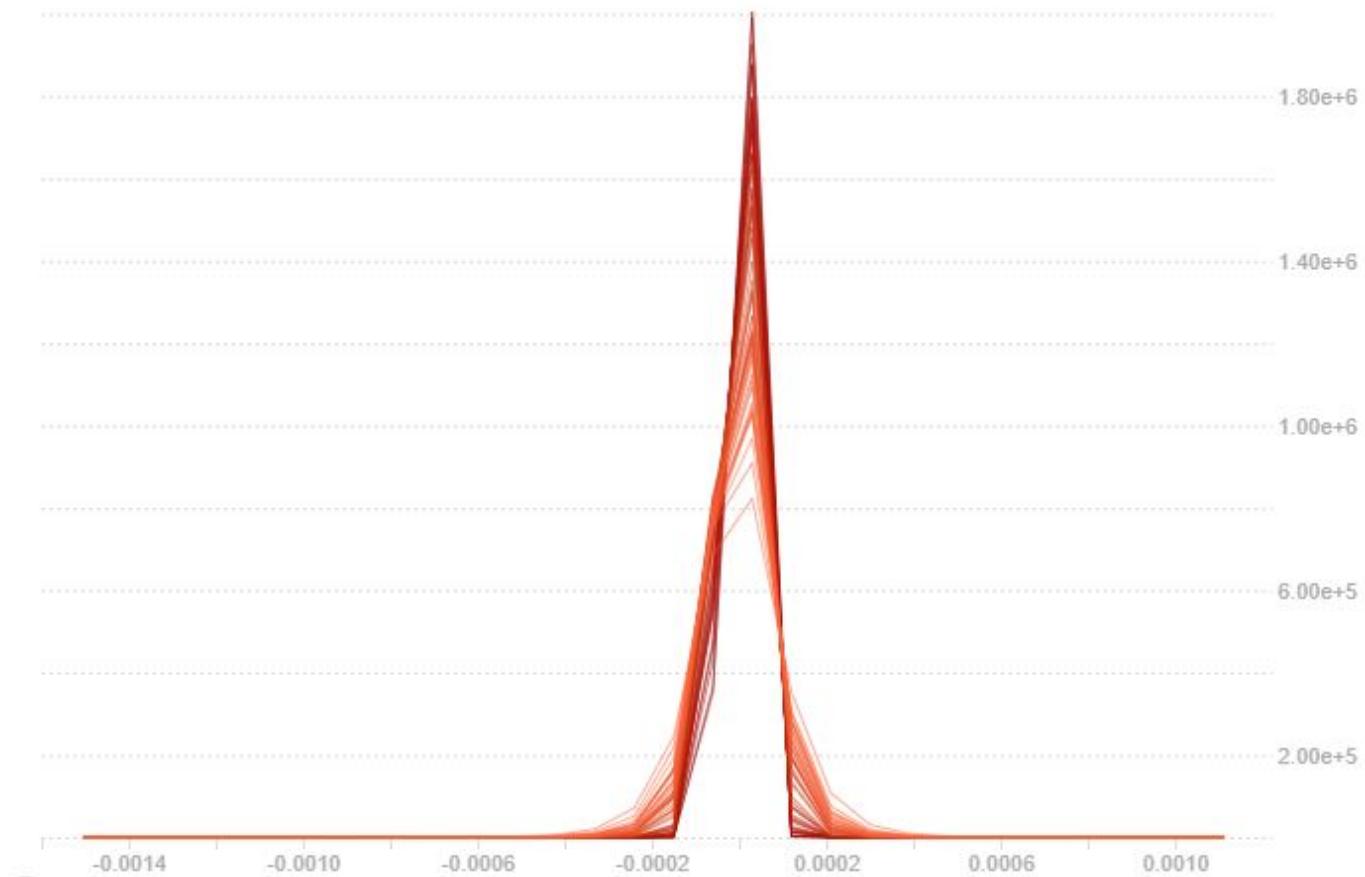
ResNet\_Q3\_histogram



The following 4 images below showcase histogram of gradients for the last two layers:

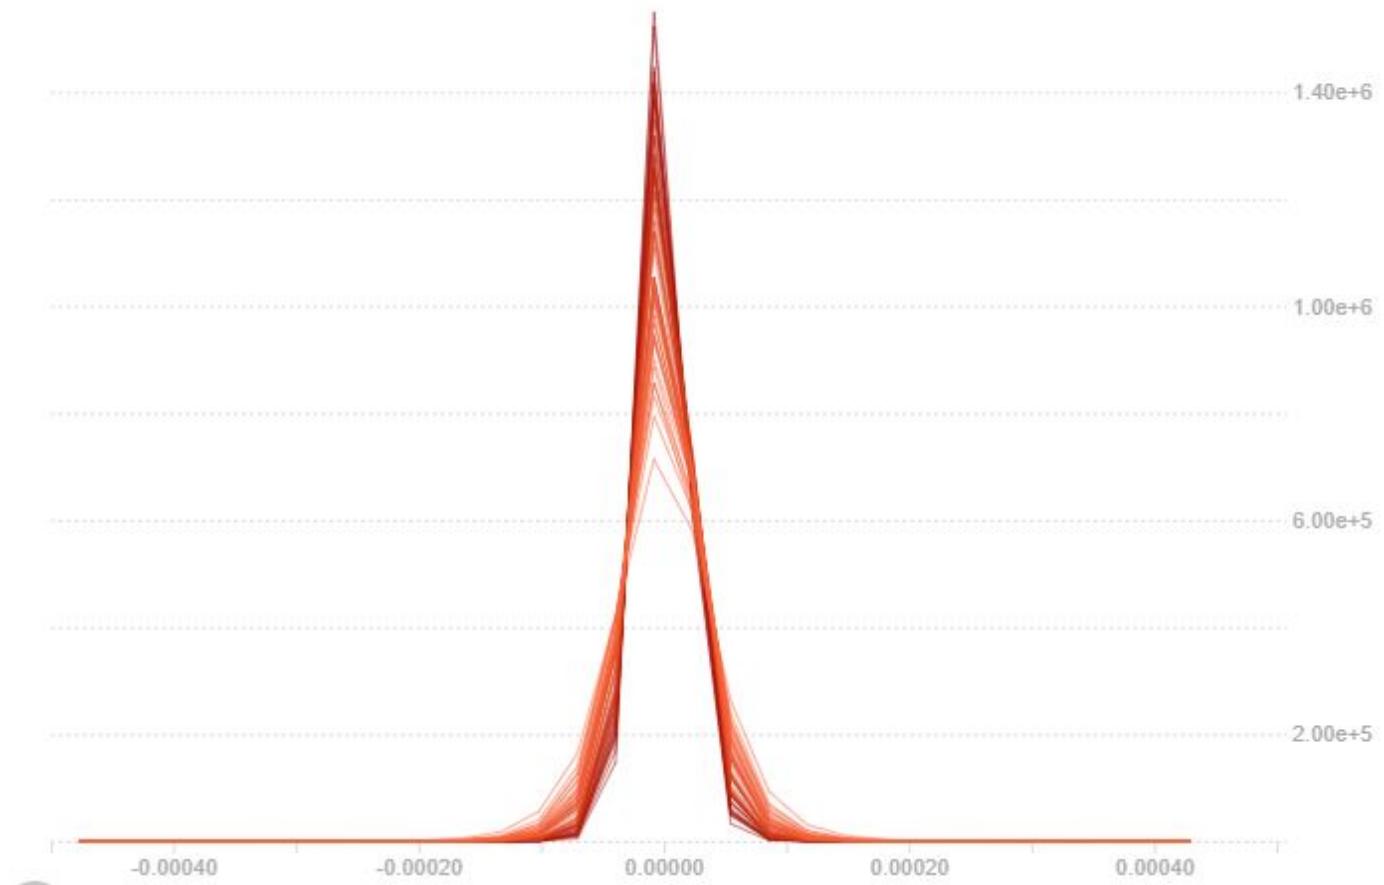
0.layer4.1.conv1.weight.grad

ResNet\_Q3\_histogram



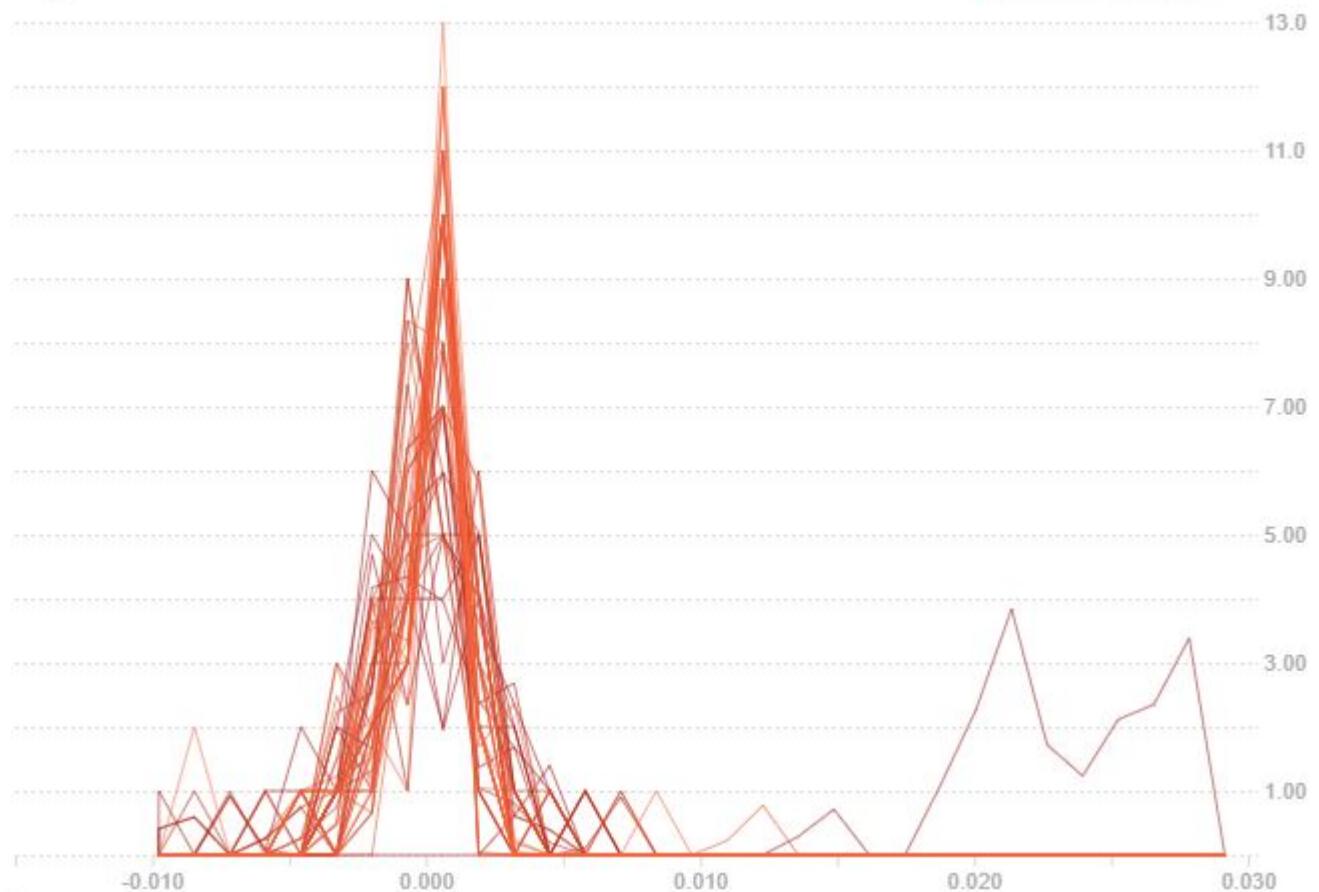
0.layer4.1.conv2.weight.grad

ResNet\_Q3\_histogram



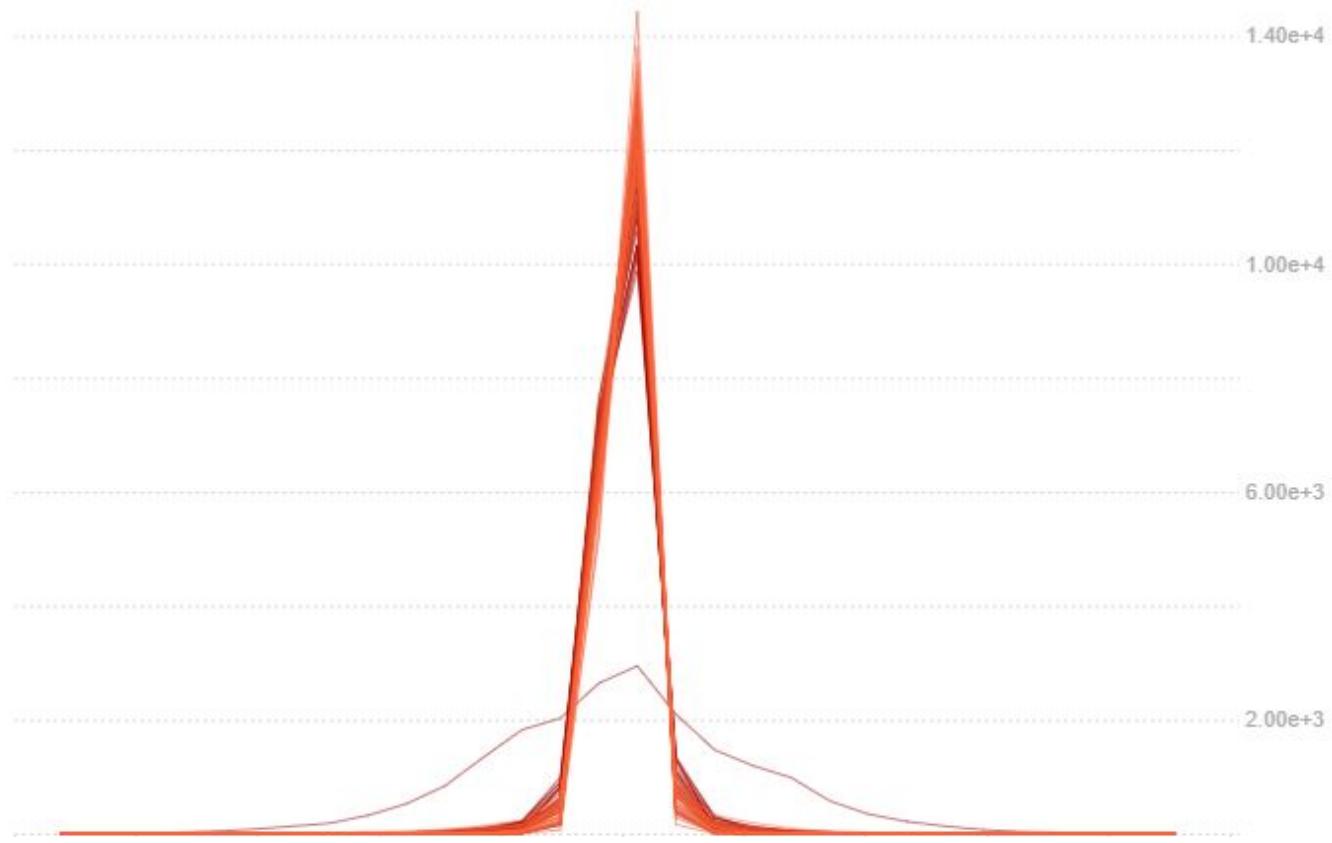
1.bias.grad

ResNet\_Q3\_histogram



1.weight.grad

ResNet\_Q3\_histogram



```
In [ ]: args = ARGs(batch_size = 32, epochs=50, lr = 0.0001)
args.gamma = 0.3
modelres = models.resnet18(pretrained=False)
model= nn.Sequential(modelres,nn.Linear(1000,20,bias=True))
optimizer = torch.optim.Adam(model.parameters(), lr = args.lr)
scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=15, gamma=args.gamma)
if __name__ == '__main__':
    test_ap, test_map = trainer.train(args, model, optimizer, scheduler)
    print('test map:', test_map)
```

## Appendix A: Training Epoch Information

The following log shown below displays the batch iteration, loss, and mAP calculation at various points during the training process:

Train Epoch: 0 [0 (0%)] Loss: 0.720259 | mAP: 0.073370  
Train Epoch: 0 [100 (64%)] Loss: 0.215044 | mAP: 0.189455  
Train Epoch: 1 [200 (27%)] Loss: 0.216712 | mAP: 0.226419  
Train Epoch: 1 [300 (91%)] Loss: 0.208612 | mAP: 0.234548  
Train Epoch: 2 [400 (55%)] Loss: 0.181140 | mAP: 0.279389  
Train Epoch: 3 [500 (18%)] Loss: 0.202289 | mAP: 0.311690  
Train Epoch: 3 [600 (82%)] Loss: 0.197823 | mAP: 0.323791  
Train Epoch: 4 [700 (46%)] Loss: 0.184834 | mAP: 0.330959  
Train Epoch: 5 [800 (10%)] Loss: 0.185951 | mAP: 0.346840  
Train Epoch: 5 [900 (73%)] Loss: 0.165875 | mAP: 0.347633  
Train Epoch: 6 [1000 (37%)] Loss: 0.165044 | mAP: 0.372633  
Train Epoch: 7 [1100 (1%)] Loss: 0.156291 | mAP: 0.397830  
Train Epoch: 7 [1200 (64%)] Loss: 0.149558 | mAP: 0.405792  
Train Epoch: 8 [1300 (28%)] Loss: 0.155647 | mAP: 0.409463  
Train Epoch: 8 [1400 (92%)] Loss: 0.146398 | mAP: 0.414038  
Train Epoch: 9 [1500 (55%)] Loss: 0.153696 | mAP: 0.396369  
Train Epoch: 10 [1600 (19%)] Loss: 0.137831 | mAP: 0.410398  
Train Epoch: 10 [1700 (83%)] Loss: 0.167851 | mAP: 0.430070  
Train Epoch: 11 [1800 (46%)] Loss: 0.130406 | mAP: 0.445091  
Train Epoch: 12 [1900 (10%)] Loss: 0.152793 | mAP: 0.440219  
Train Epoch: 12 [2000 (74%)] Loss: 0.140506 | mAP: 0.434110  
Train Epoch: 13 [2100 (38%)] Loss: 0.140575 | mAP: 0.455477  
Train Epoch: 14 [2200 (1%)] Loss: 0.118644 | mAP: 0.456327  
Train Epoch: 14 [2300 (65%)] Loss: 0.138294 | mAP: 0.445303  
Train Epoch: 15 [2400 (29%)] Loss: 0.103462 | mAP: 0.494774  
Train Epoch: 15 [2500 (92%)] Loss: 0.114315 | mAP: 0.497309  
Train Epoch: 16 [2600 (56%)] Loss: 0.105307 | mAP: 0.506209  
Train Epoch: 17 [2700 (20%)] Loss: 0.126236 | mAP: 0.505680  
Train Epoch: 17 [2800 (83%)] Loss: 0.111384 | mAP: 0.506179  
Train Epoch: 18 [2900 (47%)] Loss: 0.089492 | mAP: 0.491313  
Train Epoch: 19 [3000 (11%)] Loss: 0.129798 | mAP: 0.506378  
Train Epoch: 19 [3100 (75%)] Loss: 0.124896 | mAP: 0.506707  
Train Epoch: 20 [3200 (38%)] Loss: 0.097039 | mAP: 0.509701  
Train Epoch: 21 [3300 (2%)] Loss: 0.069134 | mAP: 0.498853  
Train Epoch: 21 [3400 (66%)] Loss: 0.109458 | mAP: 0.506017  
Train Epoch: 22 [3500 (29%)] Loss: 0.070703 | mAP: 0.500423  
Train Epoch: 22 [3600 (93%)] Loss: 0.074366 | mAP: 0.504471  
Train Epoch: 23 [3700 (57%)] Loss: 0.092572 | mAP: 0.503754  
Train Epoch: 24 [3800 (20%)] Loss: 0.100229 | mAP: 0.504639  
Train Epoch: 24 [3900 (84%)] Loss: 0.089912 | mAP: 0.508232  
Train Epoch: 25 [4000 (48%)] Loss: 0.083635 | mAP: 0.502918  
Train Epoch: 26 [4100 (11%)] Loss: 0.083434 | mAP: 0.507066  
Train Epoch: 26 [4200 (75%)] Loss: 0.070950 | mAP: 0.511265

Train Epoch: 27 [4300 (39%)] Loss: 0.073017 | mAP: 0.503205  
Train Epoch: 28 [4400 (3%)] Loss: 0.074057 | mAP: 0.509178  
Train Epoch: 28 [4500 (66%)] Loss: 0.089060 | mAP: 0.508615  
Train Epoch: 29 [4600 (30%)] Loss: 0.081091 | mAP: 0.508820  
Train Epoch: 29 [4700 (94%)] Loss: 0.082942 | mAP: 0.511548  
Train Epoch: 30 [4800 (57%)] Loss: 0.063460 | mAP: 0.518176  
Train Epoch: 31 [4900 (21%)] Loss: 0.073838 | mAP: 0.515324  
Train Epoch: 31 [5000 (85%)] Loss: 0.062788 | mAP: 0.514318  
Train Epoch: 32 [5100 (48%)] Loss: 0.067382 | mAP: 0.518381  
Train Epoch: 33 [5200 (12%)] Loss: 0.063792 | mAP: 0.510551  
Train Epoch: 33 [5300 (76%)] Loss: 0.058528 | mAP: 0.511280  
Train Epoch: 34 [5400 (39%)] Loss: 0.080845 | mAP: 0.512993  
Train Epoch: 35 [5500 (3%)] Loss: 0.071587 | mAP: 0.515748  
Train Epoch: 35 [5600 (67%)] Loss: 0.066521 | mAP: 0.516310  
Train Epoch: 36 [5700 (31%)] Loss: 0.038676 | mAP: 0.513884  
Train Epoch: 36 [5800 (94%)] Loss: 0.109776 | mAP: 0.514998  
Train Epoch: 37 [5900 (58%)] Loss: 0.056390 | mAP: 0.509673  
Train Epoch: 38 [6000 (22%)] Loss: 0.057494 | mAP: 0.517916  
Train Epoch: 38 [6100 (85%)] Loss: 0.048027 | mAP: 0.514223  
Train Epoch: 39 [6200 (49%)] Loss: 0.062421 | mAP: 0.510778  
Train Epoch: 40 [6300 (13%)] Loss: 0.064487 | mAP: 0.508616  
Train Epoch: 40 [6400 (76%)] Loss: 0.060632 | mAP: 0.512278  
Train Epoch: 41 [6500 (40%)] Loss: 0.076664 | mAP: 0.509637  
Train Epoch: 42 [6600 (4%)] Loss: 0.053598 | mAP: 0.513044  
Train Epoch: 42 [6700 (68%)] Loss: 0.054771 | mAP: 0.512391  
Train Epoch: 43 [6800 (31%)] Loss: 0.044373 | mAP: 0.511334  
Train Epoch: 43 [6900 (95%)] Loss: 0.071721 | mAP: 0.509659  
Train Epoch: 44 [7000 (59%)] Loss: 0.068732 | mAP: 0.509379  
Train Epoch: 45 [7100 (22%)] Loss: 0.108448 | mAP: 0.510053  
Train Epoch: 45 [7200 (86%)] Loss: 0.069663 | mAP: 0.512628  
Train Epoch: 46 [7300 (50%)] Loss: 0.054835 | mAP: 0.510637  
Train Epoch: 47 [7400 (13%)] Loss: 0.069490 | mAP: 0.511583  
Train Epoch: 47 [7500 (77%)] Loss: 0.059993 | mAP: 0.508102  
Train Epoch: 48 [7600 (41%)] Loss: 0.058032 | mAP: 0.512003  
Train Epoch: 49 [7700 (4%)] Loss: 0.056961 | mAP: 0.513160  
Train Epoch: 49 [7800 (68%)] Loss: 0.047862 | mAP: 0.509374  
test map: 0.5124332392148173

# Q4 Shoulders of Giants (15 points)

As we have already seen, deep networks can sometimes be hard to optimize. Often times they heavily overfit on small training sets. Many approaches have been proposed to counter this, eg, [Krahenbuhl et al. \(ICLR'16\)](#) (<http://arxiv.org/pdf/1511.06856.pdf>), self-supervised learning, etc. However, the most effective approach remains pre-training the network on large, well-labeled supervised datasets such as ImageNet.

While training on the full ImageNet data is beyond the scope of this assignment, people have already trained many popular/standard models and released them online. In this task, we will initialize a ResNet-18 model with pre-trained ImageNet weights (from `torchvision` ), and finetune the network for PASCAL classification.

## 4.1 Load Pre-trained Model (7 pts)\

Load the pre-trained weights up to the second last layer, and initialize last weights and biases from scratch.

The model loading mechanism is based on names of the weights. It is easy to load pretrained models from `torchvision.models` , even when your model uses different names for weights. Please briefly explain how to load the weights correctly if the names do not match ([hint \(<https://discuss.pytorch.org/t/loading-weights-from-pretrained-model-with-different-module-names/11841>\)](#)).

**YOUR ANSWER HERE**

If you are loading pretrained parameters into a model and the names of the parameters do not match, you can still iterate through the layers of the model and load those weights. First, when you load your pretrained parameters, turn those parameters into a list. You will be iterating through the list later. Build a for loop to iterate through your new models parameters. At each iteration of the for loop, set the layer\_name and weights of the new model to be equal to the name and parameters of the pretrained model (which was turned into a list prior to the for loop). Iterate through the layers of the new model until the end. You should have a pretrained model-imported new model now.

```
pre_trained_model=torch.load("Path to the .pth file")
new=list(pre_trained.items())

my_model_kvpair=mymodel.state_dict()
count=0
for key,value in my_model_kvpair.item():
    layer_name,weights=new[count]
    mymodel_kvpair[key]=weights
    count+=1
```

```
In [ ]: import torch
import torch.nn as nn
import torch.nn.functional as F
from torchvision import models
import matplotlib.pyplot as plt
%matplotlib inline

import trainer
from utils import ARGS
from simple_cnn import SimpleCNN
from voc_dataset import VOCDataset

# Pre-trained weights up to second-to-last layer
# final layers should be initialized from scratch!
class PretrainedResNet(nn.Module):
    def __init__(self):
        super().__init__()
        # Load resnet model
        self.modelres = models.resnet18(pretrained = True)
        for params in self.modelres.parameters():
            params.requires_grad = False

        self.model= nn.Sequential(self.modelres,nn.Linear(1000,20,bias=True))

    def forward(self, x):
        return self.model(x)
```

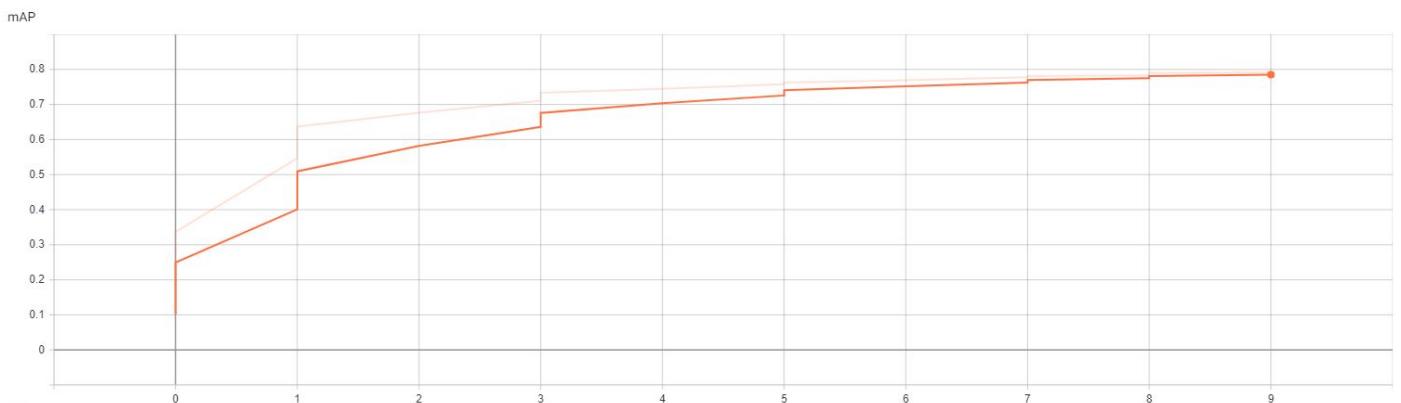
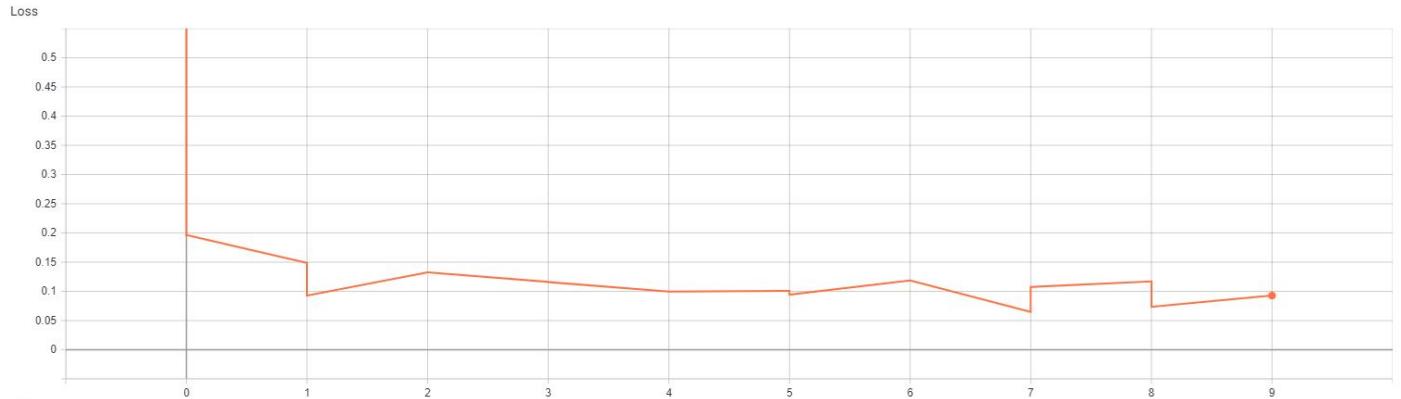
Use similar hyperparameter setup as in the scratch case. Show the learning curves (training loss, testing MAP) for 10 epochs. Please evaluate your model to calculate the MAP on the testing dataset every 100 iterations.

## REMEMBER TO SAVE MODEL AT END OF TRAINING

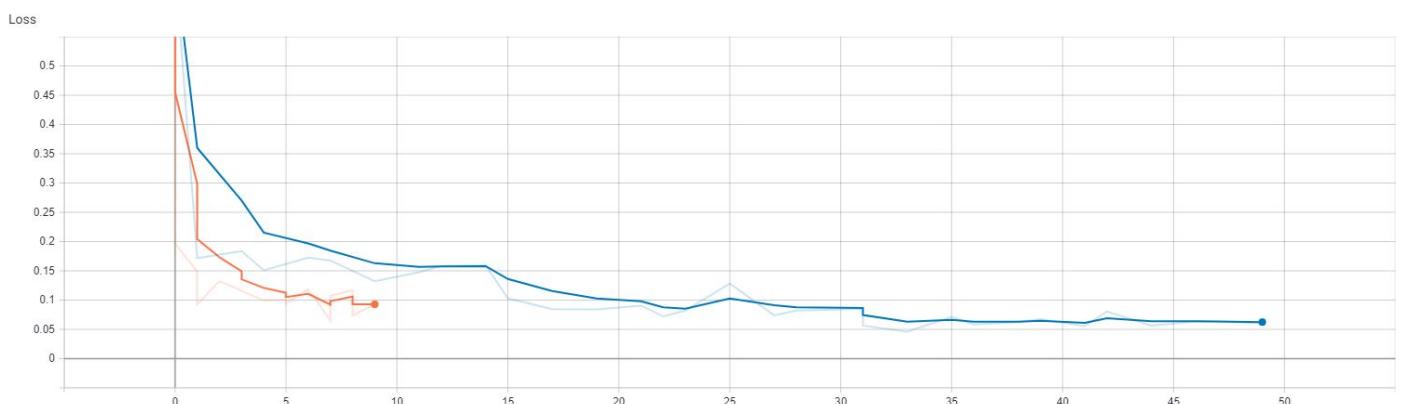
```
In [ ]: args = ARGS(batch_size = 32, epochs=10, lr = 0.0001)
args.gamma = 0.5
weightDecay = 5e-5
model = PretrainedResNet()
optimizer = torch.optim.Adam(model.parameters(), lr = args.lr, weight_decay = weightDecay)
scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=15, gamma=args.gamma)
if __name__ == '__main__':
    test_ap, test_map = trainer.train(args, model, optimizer, scheduler)
    print('test map:', test_map)
```

## YOUR TB SCREENSHOTS HERE

The following two figures shown below display the training loss and training mAP for the ResNet finetuned model.



The following two figures shown below compared the finetuned ResNet model with the ResNet model that was trained from scratch. The curves belonging to the transfer learning ResNet model are colored in light blue.





## Appendix A: Training Epoch Information

The following log shown below displays the batch iteration, loss, and mAP calculation at various points during the training process:

```
Train Epoch: 0 [0 (0%)] Loss: 0.904385 | mAP: 0.112697
Train Epoch: 0 [100 (64%)] Loss: 0.246686 | mAP: 0.318172
Train Epoch: 1 [200 (27%)] Loss: 0.145483 | mAP: 0.512955
Train Epoch: 1 [300 (91%)] Loss: 0.148025 | mAP: 0.619263
Train Epoch: 2 [400 (55%)] Loss: 0.127986 | mAP: 0.674601
Train Epoch: 3 [500 (18%)] Loss: 0.116205 | mAP: 0.710900
Train Epoch: 3 [600 (82%)] Loss: 0.118125 | mAP: 0.730045
Train Epoch: 4 [700 (46%)] Loss: 0.111438 | mAP: 0.739866
Train Epoch: 5 [800 (10%)] Loss: 0.120211 | mAP: 0.755548
Train Epoch: 5 [900 (73%)] Loss: 0.099371 | mAP: 0.761425
Train Epoch: 6 [1000 (37%)] Loss: 0.103049 | mAP: 0.767812
Train Epoch: 7 [1100 (1%)] Loss: 0.095842 | mAP: 0.778130
Train Epoch: 7 [1200 (64%)] Loss: 0.111486 | mAP: 0.776830
Train Epoch: 8 [1300 (28%)] Loss: 0.082349 | mAP: 0.783933
Train Epoch: 8 [1400 (92%)] Loss: 0.100563 | mAP: 0.788158
Train Epoch: 9 [1500 (55%)] Loss: 0.095881 | mAP: 0.792471
test map: 0.7933945926532152
```

## **Q5: Analysis (20 points)**

By now you should know how to train networks from scratch or using from pre-trained models. You should also understand the relative performance in either scenarios. Needless to say, the performance of these models is stronger than previous non-deep architectures used until 2012. However, final performance is not the only metric we care about. It is important to get some intuition of what these models are really learning. Lets try some standard techniques.

**FEEL FREE TO WRITE UTIL CODE IN ANOTHER FILE AND IMPORT IN THIS NOTEBOOK FOR EASE OF READABILITY**

### **5.1 Nearest Neighbors (7 pts)**

Pick 3 images from PASCAL test set from different classes, and compute 4 nearest neighbors of those images over the test set. You should use and compare the following feature representations for the nearest neighbors:

1. fc7 features from the ResNet (finetuned from ImageNet)
2. pool5 features from the CaffeNet (trained from scratch)

### **CaffeNet 4 Nearest Neighbors**

```
In [ ]: import torch
import torch.nn as nn
import torch.nn.functional as F
from torchvision import models
import matplotlib.pyplot as plt
%matplotlib inline

import trainer
from utils import ARGS
from simple_cnn import SimpleCNN
from voc_dataset import VOCDataset

import os
import numpy as np

#Load CaffeNet model
def get_fc(inp_dim, out_dim, non_linear='relu'):
    """
    Mid-Level API. It is useful to customize your own for Large code repo.
    :param inp_dim: int, intput dimension
    :param out_dim: int, output dimension
    :param non_linear: str, 'relu', 'softmax'
    :return: list of layers [FC(inp_dim, out_dim), (non Linear Layer)]
    """
    layers = []
    layers.append(nn.Linear(inp_dim, out_dim))
    if non_linear == 'relu':
        layers.append(nn.ReLU())
    elif non_linear == 'softmax':
        layers.append(nn.Softmax(dim=1))
    elif non_linear == 'none':
        pass
    else:
        raise NotImplementedError
    return layers

class CaffeNet(nn.Module):
    def __init__(self):
        super().__init__()
        c_dim = 3
        self.conv1 = nn.Conv2d(c_dim, 96, 11, 4, padding=0) # valid padding
        self.pool1 = nn.MaxPool2d(3,2)
        self.conv2 = nn.Conv2d(96, 256, 5, padding=2) # same padding
        self.pool2 = nn.MaxPool2d(3,2)
        self.conv3 = nn.Conv2d(256, 384, 3, padding=1) # same padding
        self.conv4 = nn.Conv2d(384, 384, 3, padding=1) # same padding
        self.conv5 = nn.Conv2d(384, 256, 3, padding=1) # same padding
        self.pool3 = nn.MaxPool2d(3,2)
        self.flat_dim = 5*5*256 # replace with the actual value
        self.fc1 = nn.Sequential(*get_fc(self.flat_dim, 4096, 'relu'))
        self.dropout1 = nn.Dropout(p=0.5)
        self.fc2 = nn.Sequential(*get_fc(4096, 4096, 'relu'))
        self.dropout2 = nn.Dropout(p=0.5)
        self.fc3 = nn.Sequential(*get_fc(4096, 20, 'none'))

        self.nonlinear = lambda x: torch.clamp(x,0)
```

```

def forward(self, x):
    N = x.size(0)
    x = self.conv1(x)
    x = self.nonlinear(x)
    x = self.pool1(x)

    x = self.conv2(x)
    x = self.nonlinear(x)
    x = self.pool2(x)

    x = self.conv3(x)
    x = self.nonlinear(x)
    x = self.conv4(x)
    x = self.nonlinear(x)
    x = self.conv5(x)
    x = self.nonlinear(x)
    x = self.pool3(x)
    x = x.view(N, self.flat_dim) # flatten the array

    out = self.fc1(x)
    out = self.nonlinear(out)
    out = self.dropout1(out)
    out = self.fc2(out)
    out = self.nonlinear(out)
    out = self.dropout2(out)
    out = self.fc3(out)

    return x

```

```

In [ ]: import utils
from sklearn.neighbors import KNeighborsClassifier

args = ARGS(batch_size = 32, use_cuda = True)

if __name__ == "__main__":
    modelCaffe = CaffeNet()
    torchLoadCaffe = torch.load('saved_models/CaffeNet-50.pth')
    modelCaffe.load_state_dict(torchLoadCaffe['model_state_dict'])
    modelCaffe = modelCaffe.to(args.device)
    modelCaffe.eval()
    testfinindex, testOutput, testTarget = trainer.train_output_CaffeNet(modelCaffe, args)

    # build k- nearest neighbors
    clf = KNeighborsClassifier(n_neighbors = 4)
    clf.fit(testOutput, testfinindex)

```

```
In [ ]: vocTest = VOCDataset(split='test',size=64)
# how big are train and test sets?
print("Test Set is: ", testOutput.shape)

# select three random images from the test set
num = 3
indexArr = []
sampleTestOutputArr = np.ones((1,6400))
sampleTestIndexArr = np.ones((1,1))

for i in range(num):
    randNum = int(np.random.rand()*len(testOutput))
    if randNum not in indexArr:
        indexArr.append(randNum)
        sampleTestOutputArr = np.concatenate((sampleTestOutputArr,testOutput[randNum][np.newaxis,:]))
        sampleTestIndexArr = np.concatenate((sampleTestIndexArr,testIndex[randNum][np.newaxis,:]))

sampleTestOutputArr = sampleTestOutputArr[1:,:]
sampleTestIndexArr = sampleTestIndexArr[1:,:]

print("Size of Test Sample Output is: ", sampleTestOutputArr.shape)
print("Size of Test Sample Output is: ", sampleTestIndexArr.shape)
```

```
In [ ]: # choose four values between 1 and 5011
testPred = clf.kneighbors(sampleTestOutputArr, return_distance=False)
print("The input is: ", np.transpose(sampleTestIndexArr))
# print("The prediction is: ",testPred)

numVal = 3

neighborImageIndex = np.zeros((numVal,4))

for i in range(numVal):
    for j in range(4):
        neighborImageIndex[i,j] = vocTest.index_list[testPred[i,j]]

print("The prediction is: ", neighborImageIndex)
```

Sample prediction output:

```
The input is: [['008515' '008131' '006491']]  
The prediction is: [[8515. 2016. 562. 3659.]  
[8131. 8356. 3707. 3514.]  
[6491. 1089. 5673. 5048.]]
```

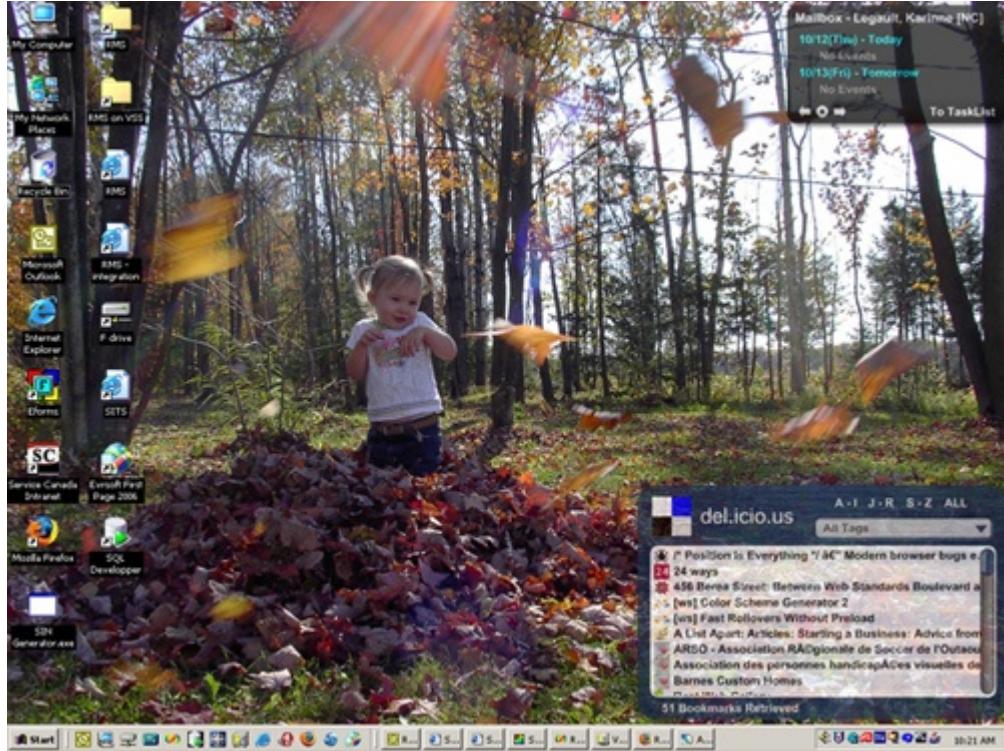
For the input image shown below ('008515.jpg'):



The follow 4 neighboring images were given:



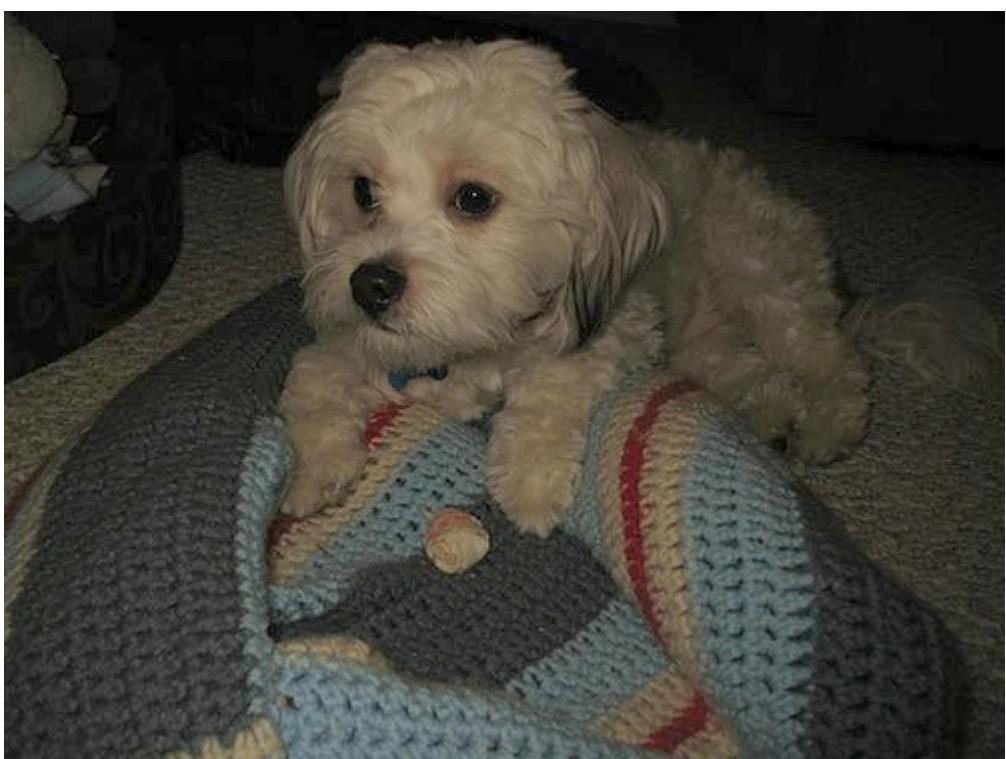




For the input image shown below ('008131.jpg'):



The follow 4 neighboring images were given:





For the input image shown below ('006491.jpg'):



The follow 4 neighboring images were given:





# ResNet 4 Nearest Neighbors Section

```
In [ ]: import torch
import torch.nn as nn
import torch.nn.functional as F
from torchvision import models
import matplotlib.pyplot as plt
%matplotlib inline

import trainer
from utils import ARGS
from simple_cnn import SimpleCNN
from voc_dataset import VOCDataset

#
import os
import numpy as np
```

```
In [ ]: # Pre-trained weights up to second-to-last Layer
# final layers should be initialized from scratch!
class PretrainedResNet(nn.Module):
    def __init__(self):
        super().__init__()
        # Load resnet model
        self.modelres = models.resnet18(pretrained = True)
        for params in self.modelres.parameters():
            params.requires_grad = False

        self.model= nn.Sequential(self.modelres,nn.Linear(1000,20,bias=True))

    def forward(self, x):
        return self.model(x)
```

```
In [ ]: import utils
from sklearn.neighbors import KNeighborsClassifier

args = ARGS(batch_size = 16, use_cuda = True)

# Create ResNet model
if __name__ == '__main__':
    m = PretrainedResNet()
    model = m.model
    model = torch.load('saved_models/PreTrainedResNet-Model.pth')
    #model.load_state_dict(torch.load('q4_resnet_pretrained_statedict.pth'))
    model = model.to(args.device)
    testindex, testOutput, testTarget = trainer.train_output_ResNet(model, args)

    clf = KNeighborsClassifier(n_neighbors = 4)
    clf.fit(testOutput, testindex)
```

```
In [ ]: vocTest = VOCdataset(split='test',size=64)

# how big are train and test sets?
print("Test Set is: ", testOutput.shape)

# select three random images from the test set
num = 3
indexArr = []
sampleTestOutputArr = np.ones((1,512))
sampleTestindexArr = np.ones((1,1))

for i in range(num):
    randNum = int(np.random.rand()*len(testOutput))
    if randNum not in indexArr:
        indexArr.append(randNum)
        sampleTestOutputArr = np.concatenate((sampleTestOutputArr,testOutput[randNum][np.newaxis,:,:]))
        sampleTestindexArr = np.concatenate((sampleTestindexArr,testindex[randNum][np.newaxis,:]))

sampleTestOutputArr = sampleTestOutputArr[1:,:]
sampleTestindexArr = sampleTestindexArr[1:,:]

print("Size of Test Sample Output is: ", sampleTestOutputArr.shape)
print("Size of Test Sample Output is: ", sampleTestindexArr.shape)
```

```
In [ ]: # choose four values between 1 and 5011
testPred = clf.kneighbors(sampleTestOutputArr, return_distance=False)
print("The input is: ", np.transpose(sampleTestIndexArr))
# print("The prediction is: ",testPred)

numVal = 3

neighborImageIndex = np.zeros((numVal,4))

for i in range(numVal):
    for j in range(4):
        neighborImageIndex[i,j] = vocTest.index_list[testPred[i,j]]

print("The prediction is: ", neighborImageIndex)
```

Sample prediction output:

```
The input is: [['000517' '003326' '001374']]  
The prediction is: [[ 517. 6194. 205. 5196.]  
 [3326. 3246. 3977. 1000.]  
 [1374. 9521. 9632. 6422.]]
```

For the input image shown below ('000517.jpg'):



The follow 4 neighboring images were given:





For the input image shown below ('003326.jpg'):



The follow 4 neighboring images were given:





For the input image shown below ('001374.jpg'):



The follow 4 neighboring images were given:







## 5.2 t-SNE visualization of intermediate features (7pts)

We can also visualize how the feature representations specialize for different classes. Take 1000 random images from the test set of PASCAL, and extract caffenet (scratch) fc7 features from those images. Compute a 2D t-SNE projection of the features, and plot them with each feature color coded by the GT class of the corresponding image. If multiple objects are active in that image, compute the color as the "mean" color of the different classes active in that image. Legend the graph with the colors for each object class.

## CaffeNet t-SNE

```
In [ ]: import torch
import torch.nn as nn
import torch.nn.functional as F
from torchvision import models
import matplotlib.pyplot as plt
%matplotlib inline

import trainer
from utils import ARGS
from simple_cnn import SimpleCNN
from voc_dataset import VOCDataset

import os
import numpy as np

#Load CaffeNet model
def get_fc(inp_dim, out_dim, non_linear='relu'):
    """
    Mid-Level API. It is useful to customize your own for Large code repo.
    :param inp_dim: int, intput dimension
    :param out_dim: int, output dimension
    :param non_linear: str, 'relu', 'softmax'
    :return: list of layers [FC(inp_dim, out_dim), (non Linear Layer)]
    """
    layers = []
    layers.append(nn.Linear(inp_dim, out_dim))
    if non_linear == 'relu':
        layers.append(nn.ReLU())
    elif non_linear == 'softmax':
        layers.append(nn.Softmax(dim=1))
    elif non_linear == 'none':
        pass
    else:
        raise NotImplementedError
    return layers

class CaffeNet(nn.Module):
    def __init__(self):
        super().__init__()
        c_dim = 3
        self.conv1 = nn.Conv2d(c_dim, 96, 11, 4, padding=0) # valid padding
        self.pool1 = nn.MaxPool2d(3,2)
        self.conv2 = nn.Conv2d(96, 256, 5, padding=2) # same padding
        self.pool2 = nn.MaxPool2d(3,2)
        self.conv3 = nn.Conv2d(256, 384, 3, padding=1) # same padding
        self.conv4 = nn.Conv2d(384, 384, 3, padding=1) # same padding
        self.conv5 = nn.Conv2d(384, 256, 3, padding=1) # same padding
        self.pool3 = nn.MaxPool2d(3,2)
        self.flat_dim = 5*5*256 # replace with the actual value
        self.fc1 = nn.Sequential(*get_fc(self.flat_dim, 4096, 'relu'))
        self.dropout1 = nn.Dropout(p=0.5)
        self.fc2 = nn.Sequential(*get_fc(4096, 4096, 'relu'))
        self.dropout2 = nn.Dropout(p=0.5)
        self.fc3 = nn.Sequential(*get_fc(4096, 20, 'none'))

        self.nonlinear = lambda x: torch.clamp(x,0)
```

```

def forward(self, x):
    N = x.size(0)
    x = self.conv1(x)
    x = self.nonlinear(x)
    x = self.pool1(x)

    x = self.conv2(x)
    x = self.nonlinear(x)
    x = self.pool2(x)

    x = self.conv3(x)
    x = self.nonlinear(x)
    x = self.conv4(x)
    x = self.nonlinear(x)
    x = self.conv5(x)
    x = self.nonlinear(x)
    x = self.pool3(x)
    x = x.view(N, self.flat_dim) # flatten the array

    out = self.fc1(x)
    out = self.nonlinear(out)
    out = self.dropout1(out)
    out = self.fc2(out)
    out1 = out
    out = self.nonlinear(out)
    out = self.dropout2(out)
    out = self.fc3(out)

    return out1

```

```

In [ ]: if __name__ == '__main__':
    modelCaffe = CaffeNet()
    torchLoadCaffe = torch.load('saved_models/CaffeNet-50.pth')
    modelCaffe.load_state_dict(torchLoadCaffe['model_state_dict'])
    modelCaffe = modelCaffe.to(args.device)
    modelCaffe.eval()
    testIndex, testOutput, testTarget = trainer.train_output_CaffeNet(modelCaffe, args)

```

```

In [ ]: # define "color" averaging function
def avg_color_label(labels):
    oneHotSum = np.sum(labels, axis=1)
    indexHotSum = np.sum(labels * np.arange(1, 21)[np.newaxis, :], axis=1)

    returnVect = indexHotSum / oneHotSum

    return returnVect[:, np.newaxis]

```

```
In [ ]: import numpy as np
from sklearn.manifold import TSNE

# select 1000 random images from the test set
num = 1000
indexArr = []
sneTestOutputArr = np.ones((1,6400))
sneTestTargetArr = np.ones((1,20))

while len(indexArr) < 1000:
    randNum = int(np.random.rand()*len(testOutput))
    if randNum not in indexArr:
        indexArr.append(randNum)
        sneTestOutputArr = np.concatenate((sneTestOutputArr,testOutput[randNum][np.newaxis,:,:]))
        sneTestTargetArr = np.concatenate((sneTestTargetArr,testTarget[randNum][np.newaxis,:,:]))
    i += 1

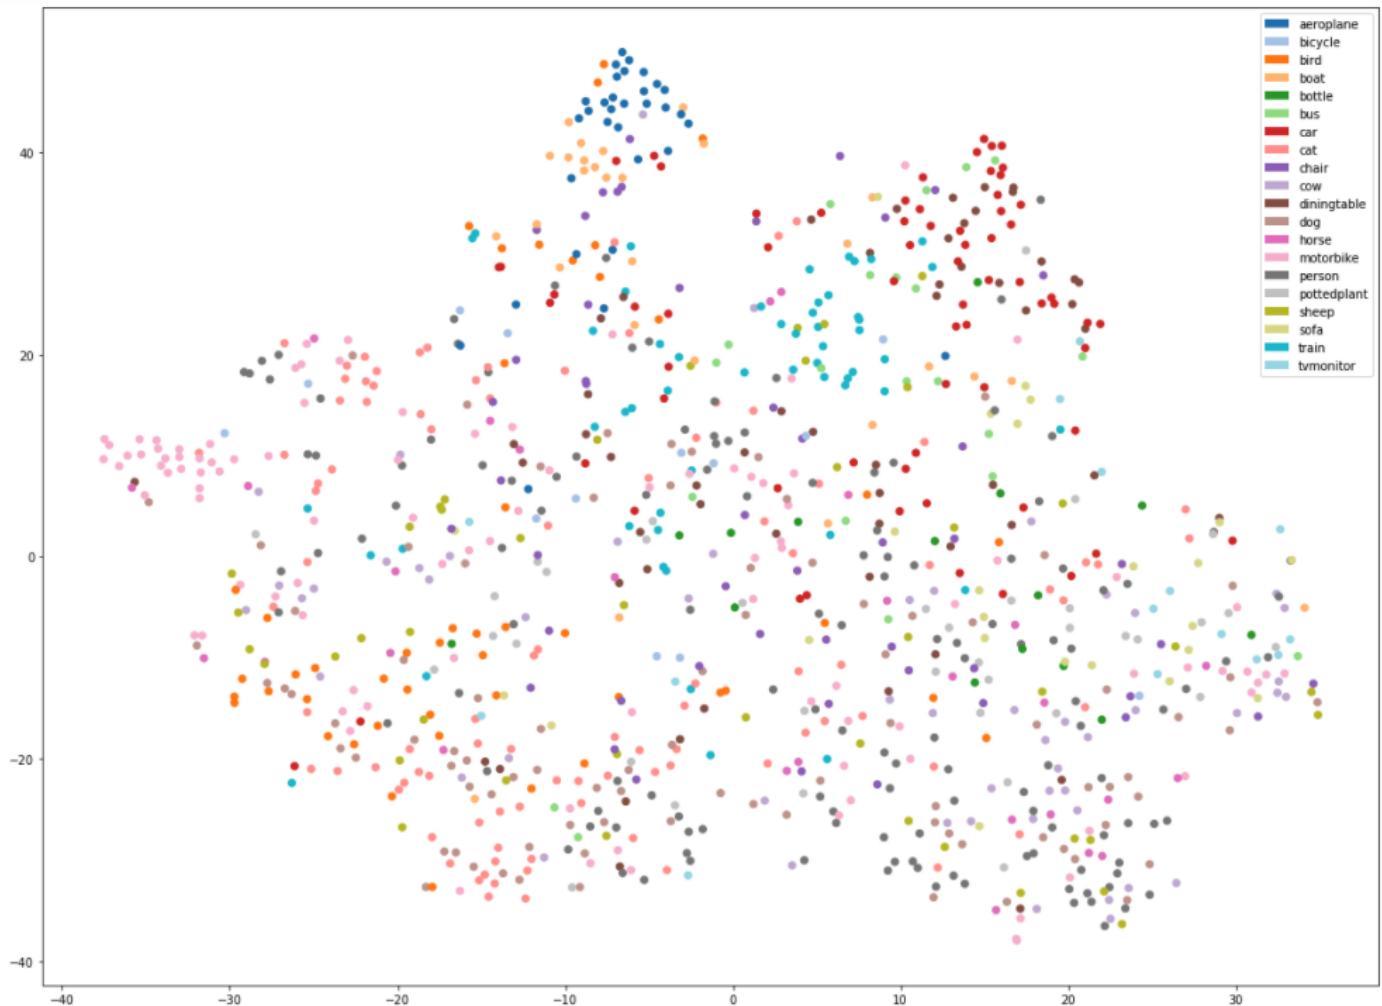
sneTestOutputArr = sneTestOutputArr[1:,:]
sneTestTargetArr = sneTestTargetArr[1:,:]

x_inbedded = TSNE(n_components=2).fit_transform(sneTestOutputArr)
```

```
In [ ]: import matplotlib

meanLabels = avg_color_label(sneTestTargetArr)
#meanLabels = meanLabels/np.amax(meanLabels)
meanLabels = np.squeeze(meanLabels, axis=1)
plt.figure(figsize=(20,15))
cmap = plt.get_cmap('tab20')
recs = []
legend = ['aeroplane', 'bicycle', 'bird', 'boat', 'bottle', 'bus', 'car',
          'cat', 'chair', 'cow', 'diningtable', 'dog', 'horse', 'motorbike',
          'person', 'pottedplant', 'sheep', 'sofa', 'train', 'tvmonitor']
plt.scatter(x_inbedded[:,0],x_inbedded[:,1],c = meanLabels, cmap = plt.get_cmap('tab20'))
for i in range(0,len(meanLabels)):
    recs.append(matplotlib.patches.Rectangle((0,0),1,1,fc=cmap(i)))
plt.legend(recs,legend,loc=1)
```

The output plot is shown as the following figure below:



## 5.3 Are some classes harder? (6pts)

Show the per-class performance of your caffenet (scratch) and ResNet (finetuned) models. Try to explain, by observing examples from the dataset, why some classes are harder or easier than the others (consider the easiest and hardest class). Do some classes see large gains due to pre-training? Can you explain why that might happen?

**YOUR ANSWER HERE**

### CaffeNet and PreTrained ResNet

The class that had the highest AP performance was the "person" while the class that had the lowest performance was the "bottle".

There can be many reasons as to why the "person" performs the best. Two images of the class are shown below for reference. One reason could be that the features of a person are the most distinct from the other classes in the dataset such as "cow" or "aeroplane". In addition, a nondifficult image of a person in the dataset exposes a lot of features such as the upright position, arms, legs, facial features, colored attire, which can help to differentiate it from the other images in the dataset during training.





Two images of the class "bottle" are shown below for reference. Compared to the features of a human, there are significantly less features of a bottle for the model to detect. In some cases, the background may make the profile harder to detect, as shown in the image below where a close-up picture was taken of some wine bottles in a dark room. In addition, features of bottles may be more and more difficult to detect as the object is further away from the camera. For example, the label of a bottle may be easier to extract as a feature when the camera is close-up, but when the bottle is a couple of meters away from the camera, then the label would be difficult for the model to spot.





## TreTrained ResNet vs From-Scratch ResNet

For reference, the per-class performance of the From-Scratch ResNet model is shown below:

aeroplane: 0.695663796029232  
bicycle: 0.5622480632184709  
bird: 0.41245039413080287  
boat: 0.5710383505097775  
bottle: 0.15982507088326053  
bus: 0.5238514422134894  
car: 0.7453615983966118  
cat: 0.46971015685819406  
chair: 0.4283246799656059  
cow: 0.28834160996828456  
diningtable: 0.40786228993246554  
dog: 0.3652161054748284  
horse: 0.7533615316892681  
motorbike: 0.6526196860641374  
person: 0.8447806695750432  
pottedplant: 0.2523940858888379  
sheep: 0.38764743196135676

sofa: 0.3606603351942273

train: 0.7182136028211066

tvmonitor: 0.4197671608743151

Though most classes saw some level of performance gain by using the transfer learning approach, the classes that saw the most significant gain were the cow, dog, and sheep classes. The gain percentage was about 3-5% in the pretrained model. One reason for this specific gain is attributed to that the model was getting better at differentiating between the three classes. As seen through some iteration with the 4 nearest neighbors model, cows, dogs, and sheeps were frequently predicted between each other, probably due to the similar features and shape between classes. Using a transfer learning approach allowed the model to run through another set of iterations to differentiate between the classes. Freezing the previous layers also allowed the model to focus on training on layer of the model, which proved to be successful in gaining performance compared to its From-Scratch counterpart.

## CaffeNet Class Performance Analysis

```
In [ ]: # redefine CaffeNet without the pool5 output
class CaffeNet(nn.Module):
    def __init__(self):
        super().__init__()
        c_dim = 3
        self.conv1 = nn.Conv2d(c_dim,96,11,4,padding=0) # valid padding
        self.pool1 = nn.MaxPool2d(3,2)
        self.conv2 = nn.Conv2d(96, 256, 5,padding=2) # same padding
        self.pool2 = nn.MaxPool2d(3,2)
        self.conv3 = nn.Conv2d(256,384,3,padding=1) # same padding
        self.conv4 = nn.Conv2d(384,384,3,padding=1) # same padding
        self.conv5 = nn.Conv2d(384,256,3,padding=1) # same padding
        self.pool3 = nn.MaxPool2d(3,2)
        self.flat_dim = 5*5*256 # replace with the actual value
        self.fc1 = nn.Sequential(*get_fc(self.flat_dim, 4096, 'relu'))
        self.dropout1 = nn.Dropout(p=0.5)
        self.fc2 = nn.Sequential(*get_fc(4096, 4096, 'relu'))
        self.dropout2 = nn.Dropout(p=0.5)
        self.fc3 = nn.Sequential(*get_fc(4096, 20, 'none'))

        self.nonlinear = lambda x: torch.clamp(x,0)

    def forward(self, x):
        N = x.size(0)
        x = self.conv1(x)
        x = self.nonlinear(x)
        x = self.pool1(x)

        x = self.conv2(x)
        x = self.nonlinear(x)
        x = self.pool2(x)

        x = self.conv3(x)
        x = self.nonlinear(x)
        x = self.conv4(x)
        x = self.nonlinear(x)
        x = self.conv5(x)
        x = self.nonlinear(x)
        x = self.pool3(x)
        x = x.view(N, self.flat_dim) # flatten the array

        out = self.fc1(x)
        out = self.nonlinear(out)
        out = self.dropout1(out)
        out = self.fc2(out)
        out = self.nonlinear(out)
        out = self.dropout2(out)
        out = self.fc3(out)

    return out
```

```
In [ ]: # output performance for each class
```

```
if __name__ == '__main__':
    modelCaffe = CaffeNet()
    torchLoadCaffe = torch.load('saved_models/CaffeNet-50.pth')
    modelCaffe.load_state_dict(torchLoadCaffe['model_state_dict'])
    modelCaffe = modelCaffe.to(args.device)
    modelCaffe.eval()

    test_loader = utils.get_data_loader('voc', train=False, batch_size=args.test_batch_size, split='test')
    ap, map = utils.eval_dataset_map(modelCaffe, args.device, test_loader)

    # output values from CaffeNet, compare with the target values
    classNames = ['aeroplane', 'bicycle', 'bird', 'boat', 'bottle', 'bus', 'car',
                  'cat', 'chair', 'cow', 'diningtable', 'dog', 'horse',
                  'motorbike', 'person', 'pottedplant', 'sheep', 'sofa', 'train', 'tvmonitor']

    print("Accuracy Precision of CaffeNet Among Individual Classes: ")
    for i in range(len(classNames)):
        print("{}: {}".format(classNames[i],ap[i]))
```

Accuracy Precision of CaffeNet Among Individual Classes:

aeroplane: 0.6581140150639725  
bicycle: 0.4561887044476487  
bird: 0.35825532448861536  
boat: 0.5089389243663608  
bottle: 0.16395719924468152  
bus: 0.3739061674684568  
car: 0.6706960925944321  
cat: 0.38321888683805233  
chair: 0.39942472719290306  
cow: 0.2317710070030363  
diningtable: 0.29885085506558173  
dog: 0.3274386268412313  
horse: 0.6648294880873716  
motorbike: 0.5684288971068936  
person: 0.8086725636680325  
pottedplant: 0.23318258506171002  
sheep: 0.25812258618853845  
sofa: 0.3496855596933779  
train: 0.5634974187463156  
tvmonitor: 0.36742932492025593

mAP:

0.4322304477043734

# ResNet Class Performance Analysis

```
In [ ]: if __name__ == '__main__':
    modelRes.eval()
    test_loader = utils.get_data_loader('voc', train=False, batch_size=args.test_batch_size, split='test')
    ap, map = utils.eval_dataset_map(modelRes, args.device, test_loader)

    # output values from CaffeNet, compare with the target values
    classNames = ['aeroplane', 'bicycle', 'bird', 'boat', 'bottle', 'bus', 'car',
                  'cat', 'chair', 'cow', 'diningtable', 'dog', 'horse',
                  'motorbike', 'person', 'pottedplant', 'sheep', 'sofa', 'train', 'tvmonitor']

    print("Accuracy Precision of PreTrained ResNet Among Individual Classes: ")
)
for i in range(len(classNames)):
    print("{}: {}".format(classNames[i],ap[i]))
```

Accuracy Precision of PreTrained ResNet Among Individual Classes:

aeroplane: 0.7213222117720576  
bicycle: 0.5694538600068458  
bird: 0.4367706344884813  
boat: 0.5773203950369964  
bottle: 0.1818182675529317  
bus: 0.5025989148818082  
car: 0.7382160645945922  
cat: 0.4581863265778226  
chair: 0.43791839747272837  
cow: 0.32130702694874635  
diningtable: 0.4435263229389825  
dog: 0.4043789661300173  
horse: 0.7495443845590672  
motorbike: 0.6876729041809619  
person: 0.854338868879449  
pottedplant: 0.2733504707389632  
sheep: 0.4169879431619927  
sofa: 0.3385033253895445  
train: 0.6906721981250176  
tvmonitor: 0.46132828158342576

mAP:

0.5132607882510216